

A photograph of the Austin, Texas skyline at dusk. In the foreground, the Congress Avenue Bridge spans the Colorado River. A massive swarm of bats is captured in flight against the darkening sky above the bridge. The city skyline, featuring the Frost Bank Tower and other skyscrapers, is visible in the background.

Shiny in Production

A blue-tinted photograph of the Austin skyline at dusk or night. A large bridge arches across the frame, and numerous small white dots representing birds are scattered across the sky.

Welcome

Sean Lopp

Kelly O'Briant

Tonya Filz

Mark Sellors

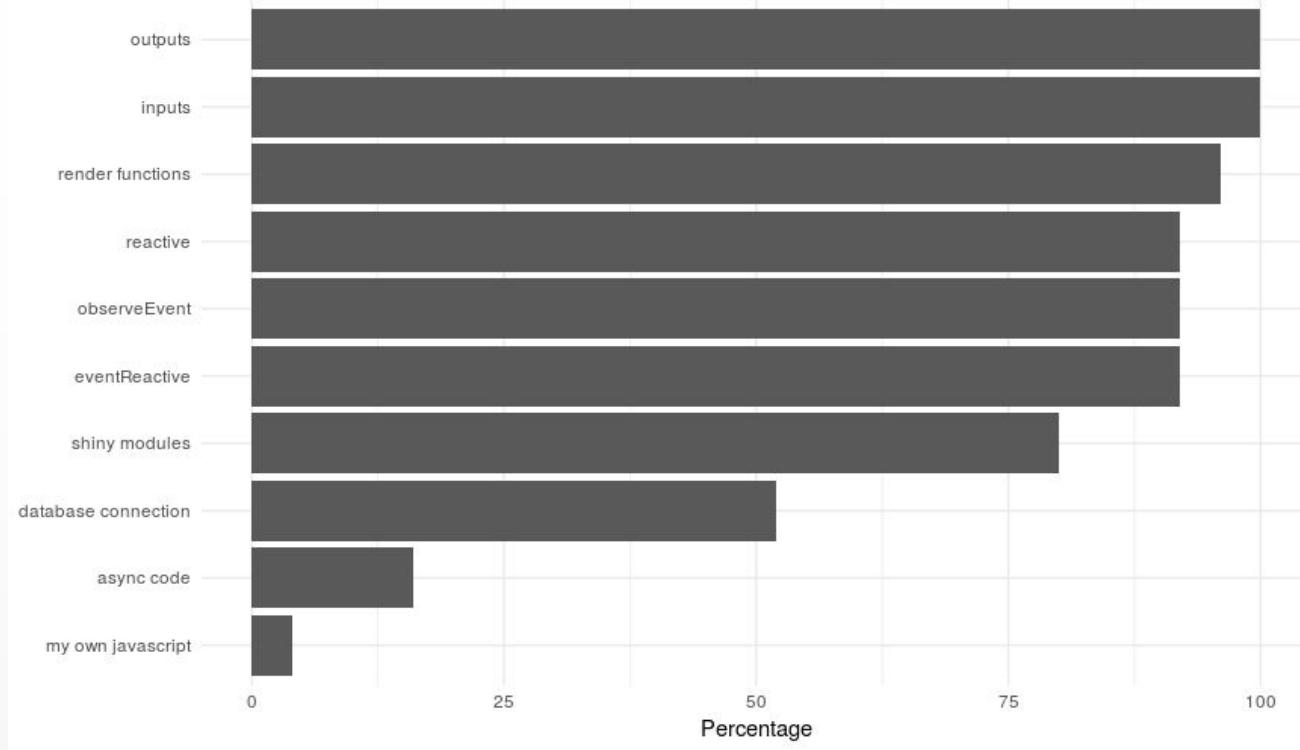
Devin Pastoor

Alan Dipert

Barret Schloerke

Survey Results

I've use ____ in my shiny app...



Objectives

Running Shiny in Production

- Is Shiny ready for production?
- What does it take to get there?
- What do you need to know?
- What tools are available?
- How do you develop a workflow?

Meet our App

- Awesome App Origin Story

TL;DR:

- We've developed a nice app
- We want to put it in production
- We want confidence that it will perform well in production, both now and in the future

App Exploration 15 Minutes

First: Open app.R and Run the Application

Discussion: *Explore the Application*

- Are there any parts of the app code that don't make sense?
- Is this application ready for production?
- How would you define "production"?
- What insights would be useful to have about the application before we try to deploy it?

Deliverable: *Start a Plan*

Create a checklist

Outline the steps you might take to put this application into production

Checklist

- Tests
- Performance Optimization
- Environment (Packages) Management
- Data Access
- Deployment Hand Off
- Scaling
- Monitoring

shinytest



Testing

Expectation Based Testing (Units Tests)

```
trim_whitespace <- function(characters) {  
  for (c in characters) {  
    if (c in whitespaces) {remove(c)}  
  }  
}  
  
expect_that(trim_whitespace("the test"), "thetest")
```

Testing

Expectation Based Testing (Units Tests)

```
trim_whitespace <- function(characters) {  
  for (c in characters) {  
    if (c in whitespaces) {remove(c)}  
  }  
}  
  
expect_that(trim_whitespace("the test"), "thetest")
```



Why is this type of testing hard for Shiny apps?

Testing

Snapshot based testing

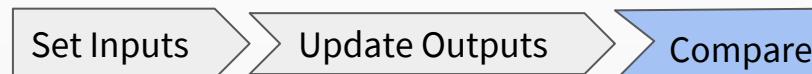
1. Take a snapshot



2. Make changes to your code



3. Compare new results to the approved snapshot



Shiny Test Demo

shinytest 20 Minutes

First: Open app.R and use “Record Test”

Discussion:

Understanding shinytest

- What does the recording file create?
- What challenges might our app pose to testing?
- Can you run the tests and get a success?
- How might you automate tests?

Deliverable: *Try shinytest*

A set of tests that can catch unintentional changes to our app.

shinytest: answers

Discuss the need for a random seed

Checkout answers/tests

Where and when should these tests run?

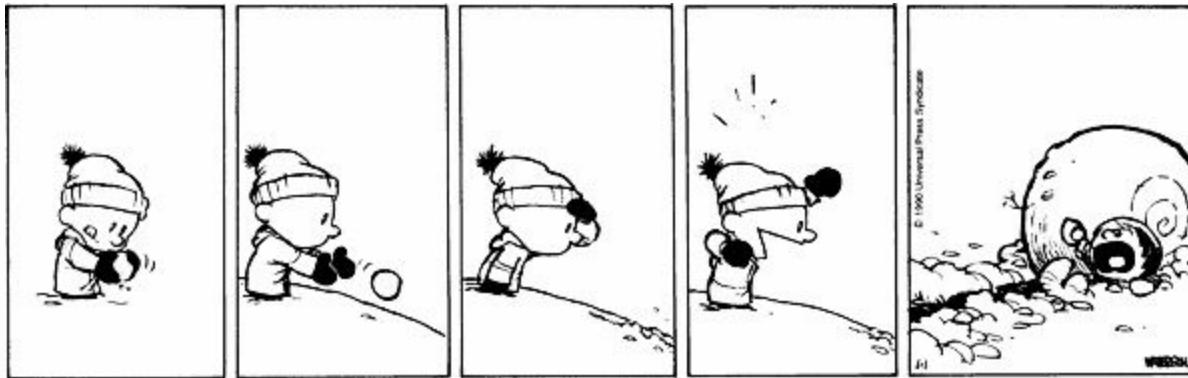
Break

Come back at 11

Profiling



The Most Important Tool



The Most Important **Tool**

THE SECOND OPTION
FEELS RIGHT. LET'S
GO WITH THAT.



Dilbert.com DilbertCartoonist@gmail.com

SHOULD WE ALWAYS
IGNORE WHAT THE
DATA SAYS, OR IS THIS
MORE OF A ONE-TIME
THING?



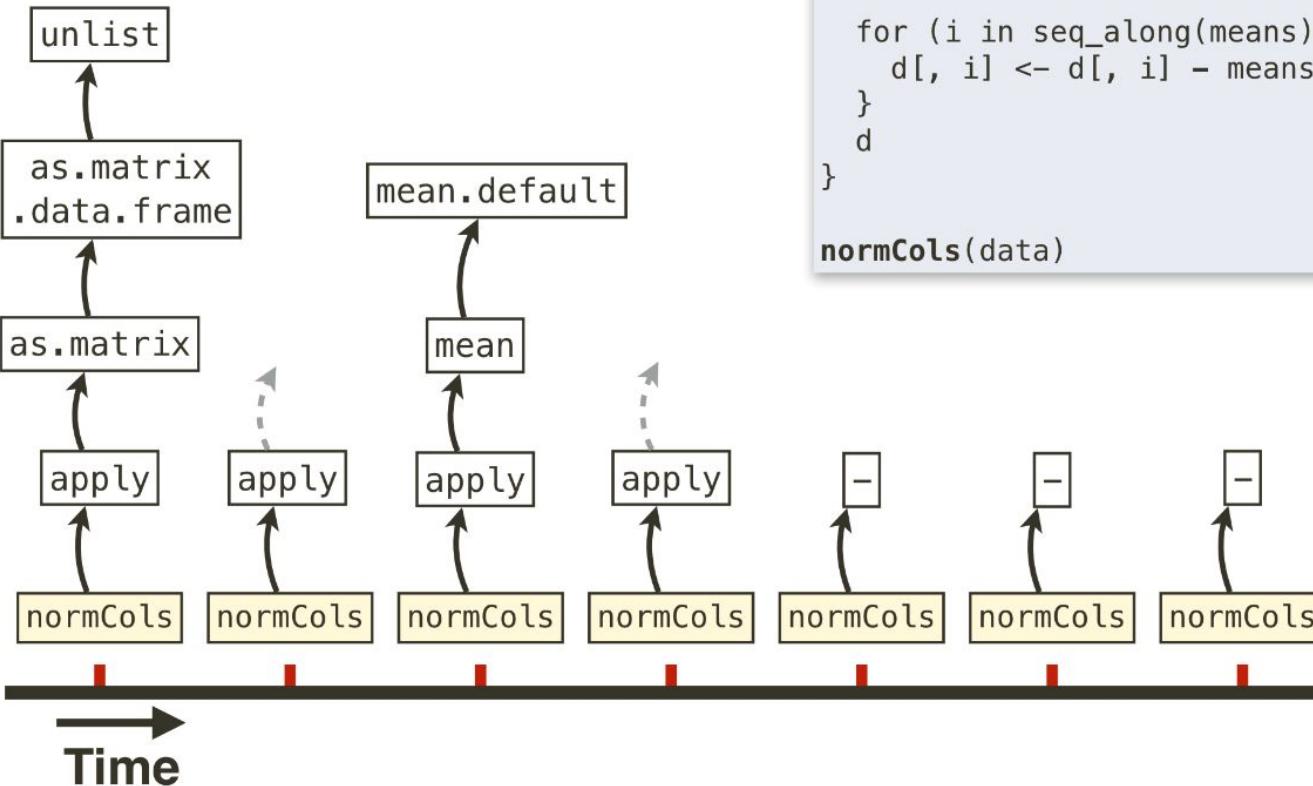
8-17-11 © 2011 Scott Adams, Inc./Dist. by Universal Uclick

IT'S
CALLED
INTUITION.



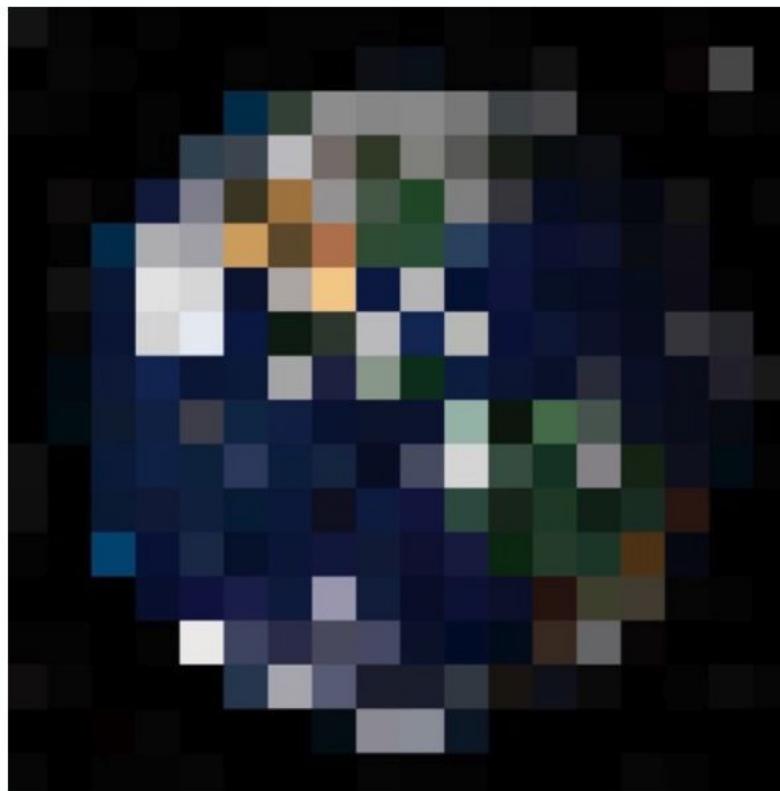
IT'S A
SLIPPERY
SLOPE TO
WITCH-
CRAFT.

Sampling profiler



```
normCols <- function(d) {  
  means <- apply(d, 2, mean)  
  
  for (i in seq_along(means)) {  
    d[, i] <- d[, i] - means[i]  
  }  
  d  
}  
  
normCols(data)
```

Profiling



profvis + shiny

```
profvis::profvis({  
  shiny::runApp()  
  # interact with app  
  # close app  
})
```

profiling 30 Minutes

First: Create a profile of our app

Discussion:

Understanding app performance

- Where does our app spend most of its time?
- Do any parts of the profile surprise you?
- If you test the app again, do you get the same results?

Deliverable: *Optimization Recommendations*

One recommendation for how we could speed up our app.

profvis + shiny: answers

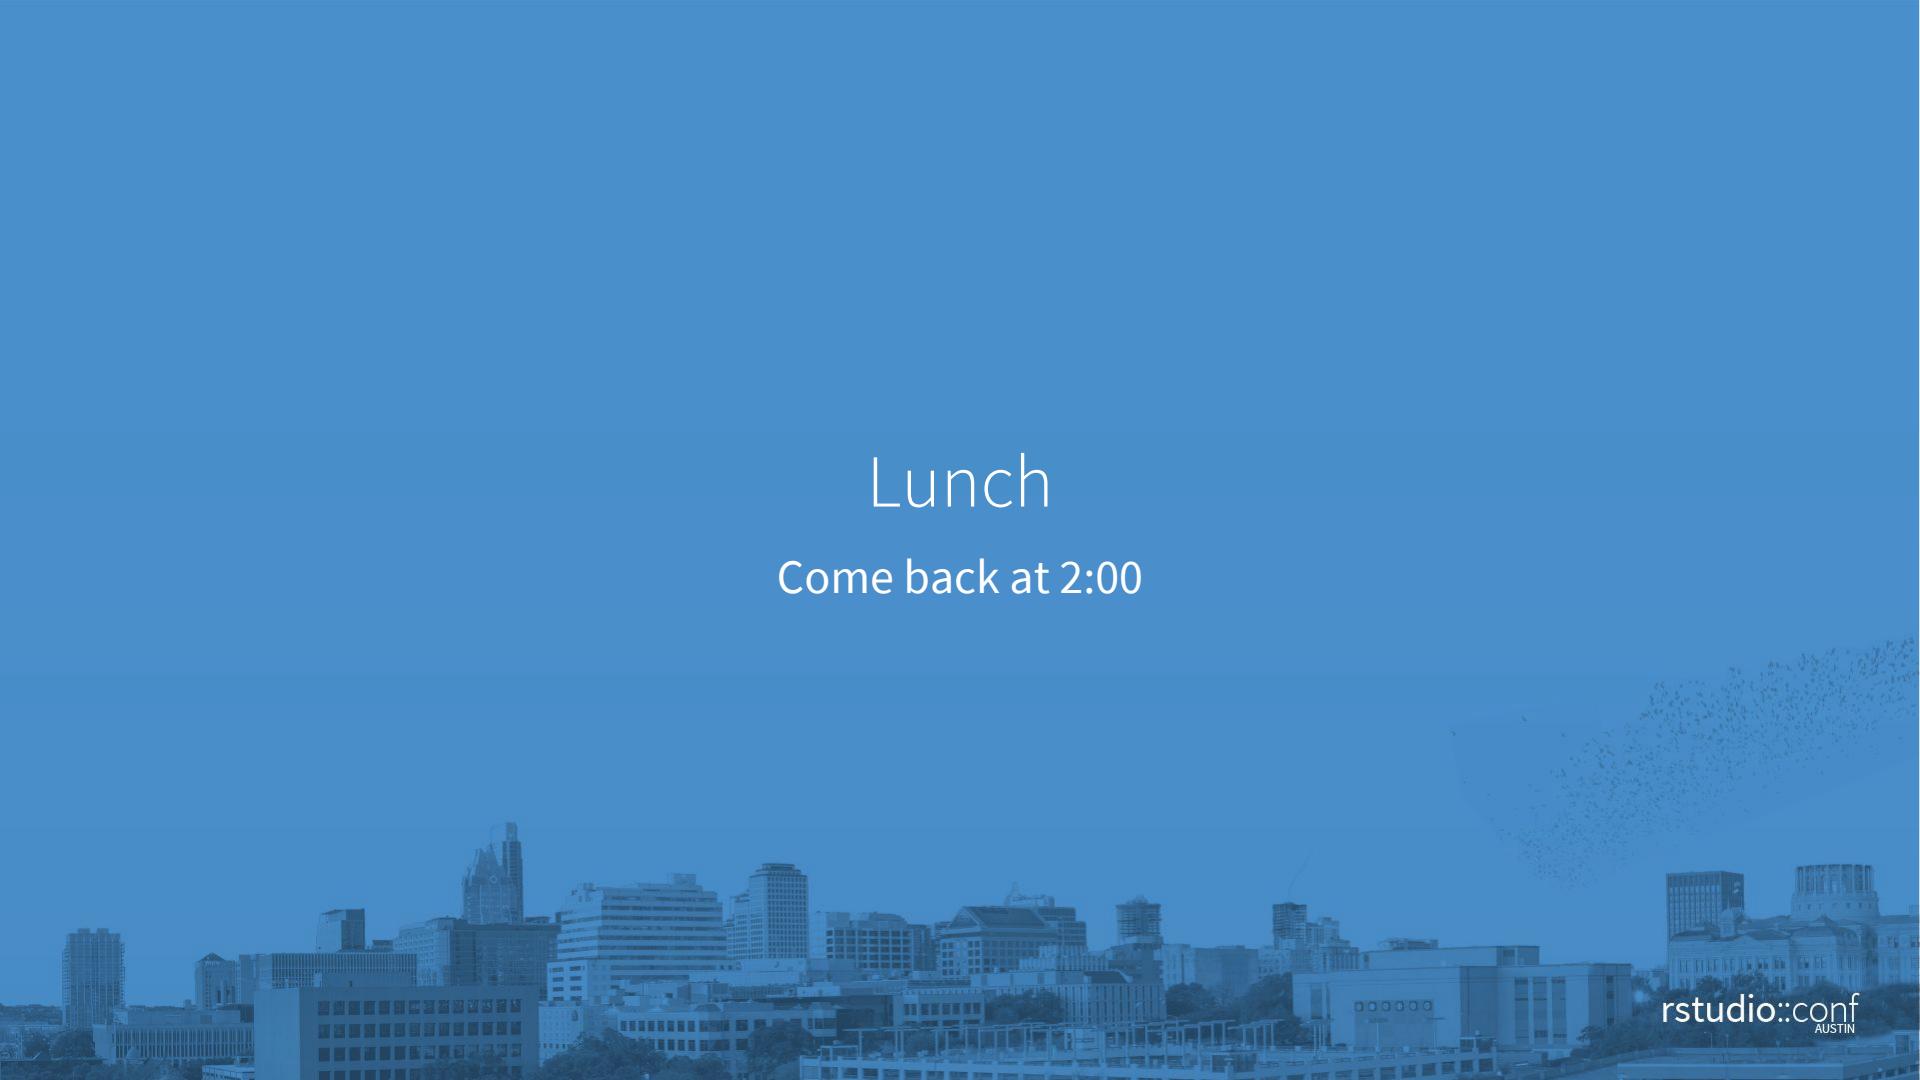
answers/app-profile.Rprofvis

Slowest thing is the lime explanation model & plotting

renderCachePlot will come to the rescue!

Other Q&A to discuss:

- Is the call to gt slow? (Answer: no, it is really fast!)
- Why does getStudentNum take so much longer than getStudentBin? How could we speed this up? (Answer: move the aggregation out of the session)

A dark blue-tinted photograph of the Austin, Texas skyline at dusk or night. The city lights are visible against a dark sky, with the tallest building, the Frost Bank Tower, clearly visible on the left.

Lunch

Come back at 2:00

Deployment



Deployment



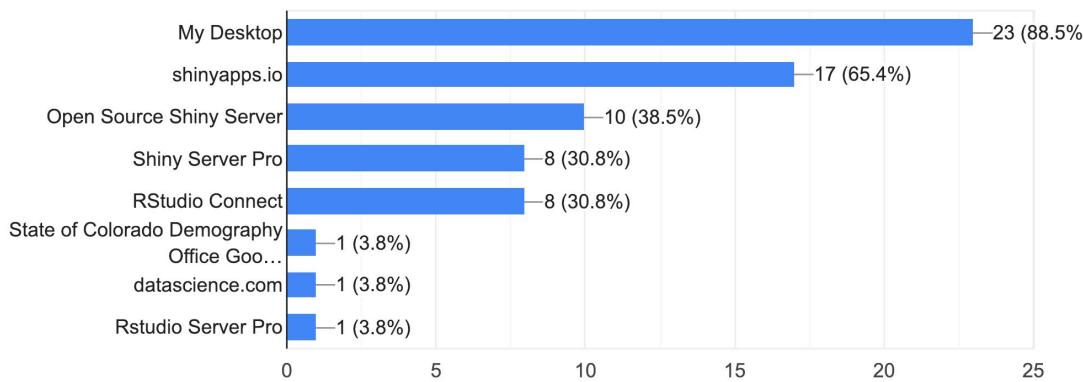
Game Plan

Do a little bit of deployment now.

Talk about scaling, deployment architectures, and other DevOps things tomorrow.

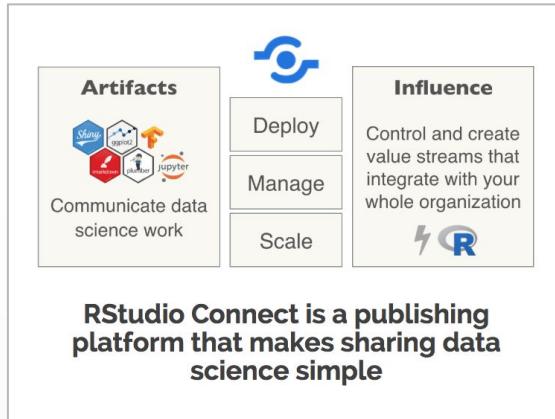
I've run a shiny application before on:

26 responses



Briefly: RStudio Connect & Why We Built It

Helping data scientists communicate their work is of utmost importance.



It doesn't matter how great your analysis is unless you can explain it to others: you need to communicate your results.

- Grolemund & Wickham in [R For Data Science](#)

A screenshot of a tweet from Amelia McNamara (@AmeliaMN). The tweet includes a graphic comparing two ways of thinking about academic goals. The top part, "How I thought of my goals in grad school:", shows a horizontal scale from "Idea" (Less valuable) to "Published paper" (More valuable), with "Preliminary results" and "Draft manuscript" in between. The bottom part, "How I should have been thinking of them:", shows a similar scale but with "Anything still on your computer" (Less valuable) at one end and "Anything out in the world" (More valuable) at the other, with "(Paper, preprint, product, blog post, open source, tweet)" in between. The tweet was posted at 11:00 AM - 3 Nov 2017 and has 340 Retweets and 855 Likes.

Deploy 15 Minutes

First: Login to RStudio Connect

<http://<your-server/rsconnect>

Username: rstudio Password: rstudio

Discussion:

Pre-deployment Brainstorm

- What is our goal in deploying this code?
- What does our code depend on locally?
- What needs to deploy with our code?

Deliverable: *Deployed App*

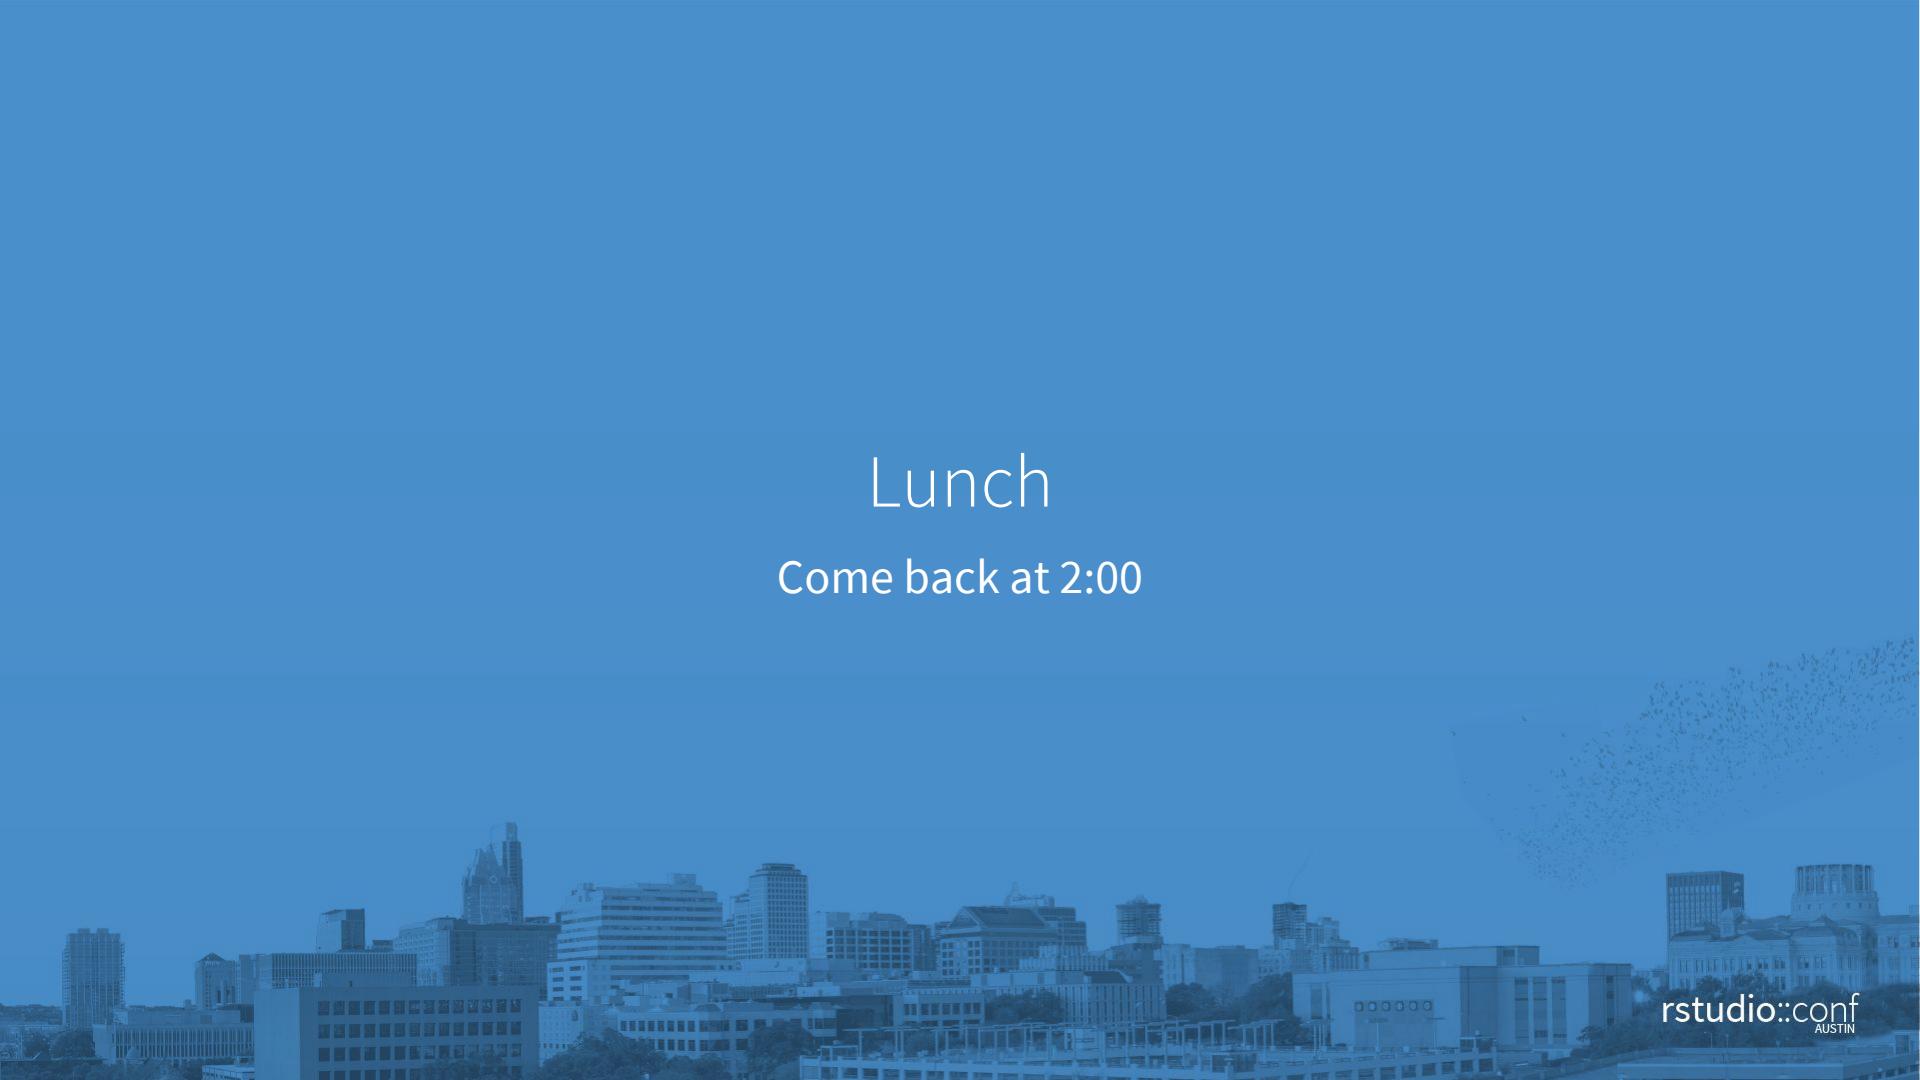
Press the publish button in RStudio:



- Link to your Connect account
- Select the files to publish (do **NOT** use the default to publish *all* the files)
- Poke around at the results

Note: We are expecting an error! That is ok

password authentication failed for user
"rstudioadmin"

A dark blue-tinted photograph of the Austin, Texas skyline at dusk or night. The city lights are visible against a dark sky, with the tallest building, the Frost Bank Tower, clearly visible on the left.

Lunch

Come back at 2:00

Deploy: answers

- Supporting Files (use_models.R, *.RDS, www/rstudio.png)
- Database Connection
 - Database Driver
 - Configuration (Credentials, DSN, config.yml, etc)
- R Version
- R Packages
- ???

Databases & config

```
library(odbc)
library(DBI)
library(config)
db <- get('database')

pool <- dbConnect(
  odbc::odbc(),
  Driver      = db$Driver,
  Server      = db$Server,
  Database    = db$Database,
  UID         = db$UID,
  PWD         = db$PWD,
  Port        = db$Port
)
```

default:

database:

Driver: PostgreSQL

Server:

'sol-eng-shiny-prod-class.cih
ykudhzbgw.us-west-2.rds.amazo
naws.com'

Database: 'students'

UID: 'rstudioadmin'

PWD: 'rstudioadmin'

Port: 5432

Databases 25 Minutes

First: Read <https://db.rstudio.com/best-practices/portable-code/>

Discussion:

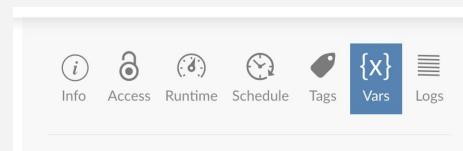
Data Management

- How does your organization connect to data?
- What data are we exposing to viewers? Is it appropriate for all viewers?
- What else could we manage with the config package?

Deliverable: *Running, Deployed App*

Update config.yml and Redeploy

Edit the Environment Variables in RStudio Connect



Refresh and



Deploy Checklist

- ☒ Supporting Files (use_models.R, *.RDS)
- ☒ Database Connection
 - ☒ Database Driver
 - ☒ Configuration (Credentials, DSN, config.yml, etc)
 - ? Think about data sensitivity & viewership

- R Version
- R Packages

packrat

```
rsconnect::writeManifest(  
  appDir = ".",
  appFiles = c("app.R", "config.yml", "model.RDS", "use_models.R",
             "model_explainer.RDS", "data_preprocessor.RDS",
             "www/rstudio.png")
)
```

Alternatives to Shiny

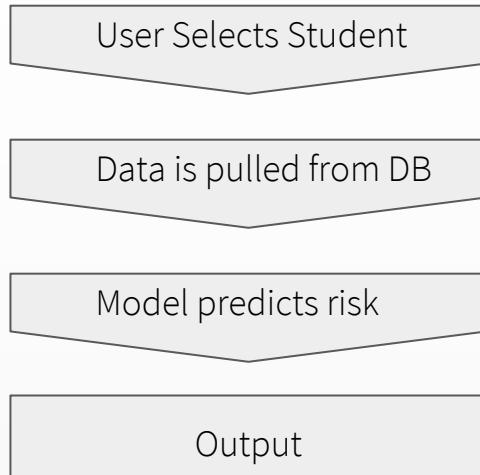
Review survey results for plumber and R Markdown familiarity



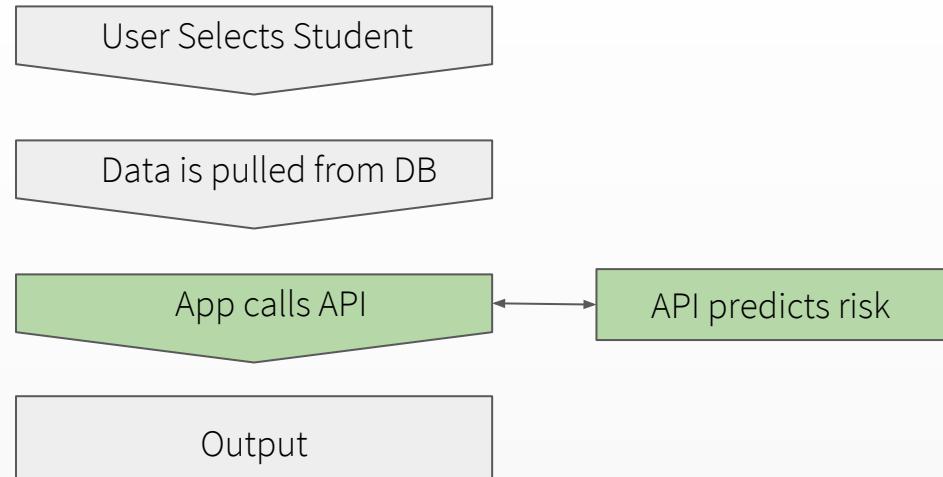
Plumber

Could our student “scoring” be hosted outside of the Shiny app?

Current Setup



Alternative Setup



Plumber 20 Minutes

First: Try creating your own plumber API in the IDE (New File)

Discussion:

Consider our new architecture

- Why might a RESTful API be easier to scale than a Shiny app?
- What benefits come from pulling the modeling code out of our app?

Deliverable:

Add an item to our production checklist to consider plumber.

Add 2 decision criteria to the checklist to decide when a plumber API would be useful

Plumber: answers

Decision Criteria:

- Does my R functionality need to be accessed by other systems?
- Do I want to be able to update parts of my code independently from the app?

Why might a RESTful API be easier to scale than a Shiny app?

Requests are stateless, so it is easier for R processes to come and go

Break

Come back at 4:00

Joe Cheng War Stories & QA



Wrap Up

First: Brainstorm how you could apply what we've learned today

Discussion:

- Share your brainstorming ideas

Deliverable A:

- Write down 3 concerns or questions you have for tomorrow.

Deliverable B:

- Do you want to present your own 5 minute “war story” tomorrow morning?
Come talk to us.

A photograph of the Austin, Texas skyline at dusk. In the foreground, the Congress Avenue Bridge spans the Colorado River. A massive swarm of bats is captured in flight against the darkening sky above the bridge. The city skyline in the background features several prominent skyscrapers, including the Frost Bank Tower and the W Hotel. The overall atmosphere is moody and atmospheric.

Day 2

War Stories

Alan Dipert (RStudio)

Devin Pastoor (Metrum)

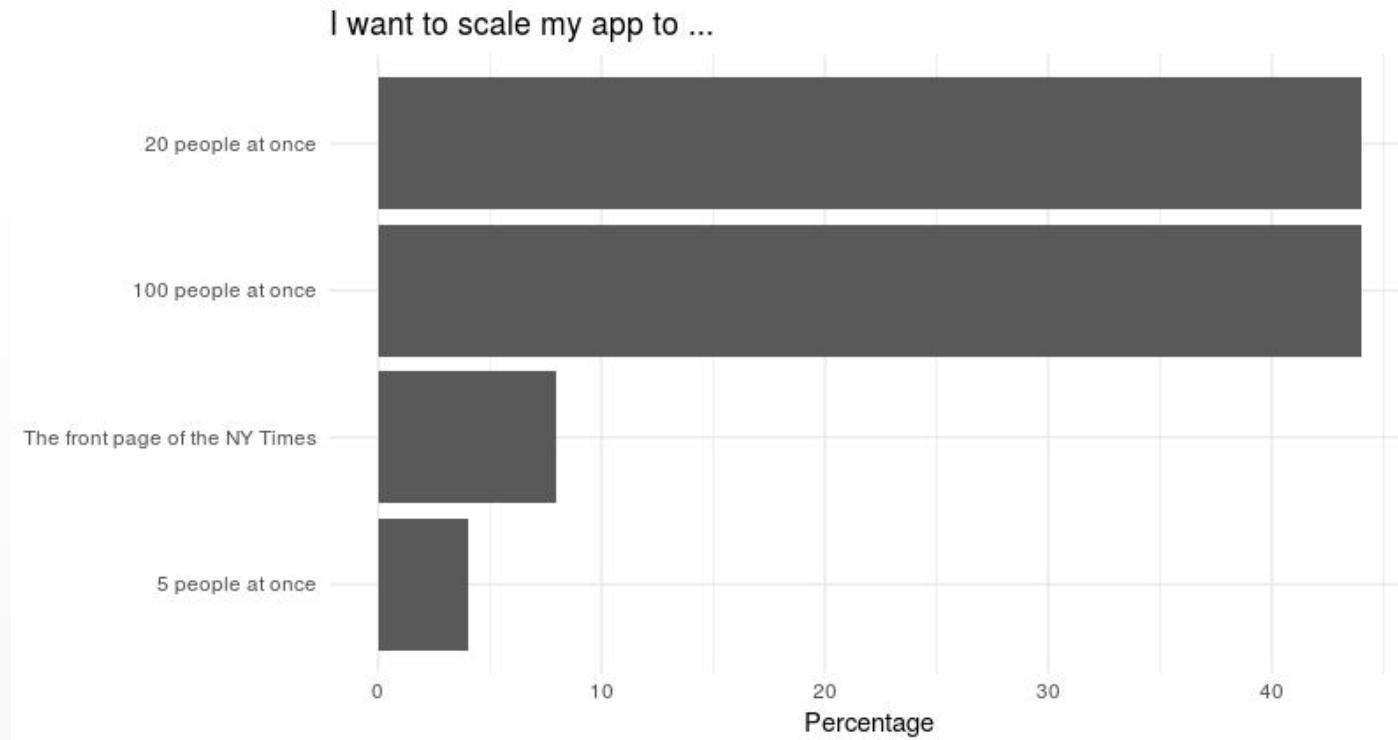
Load Testing



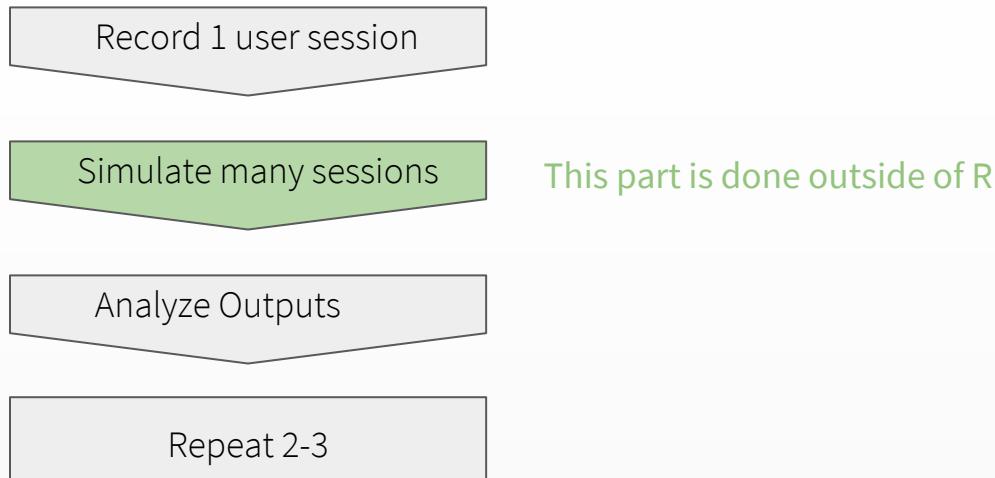
shinyloadtest



shinyloadtest



shinyloadtest



This part is done outside of R

shinyloadtest 20 Minutes

First: Open `runloadtest.R`, do the pre-run checklist

Discussion:

Preparing for the load test

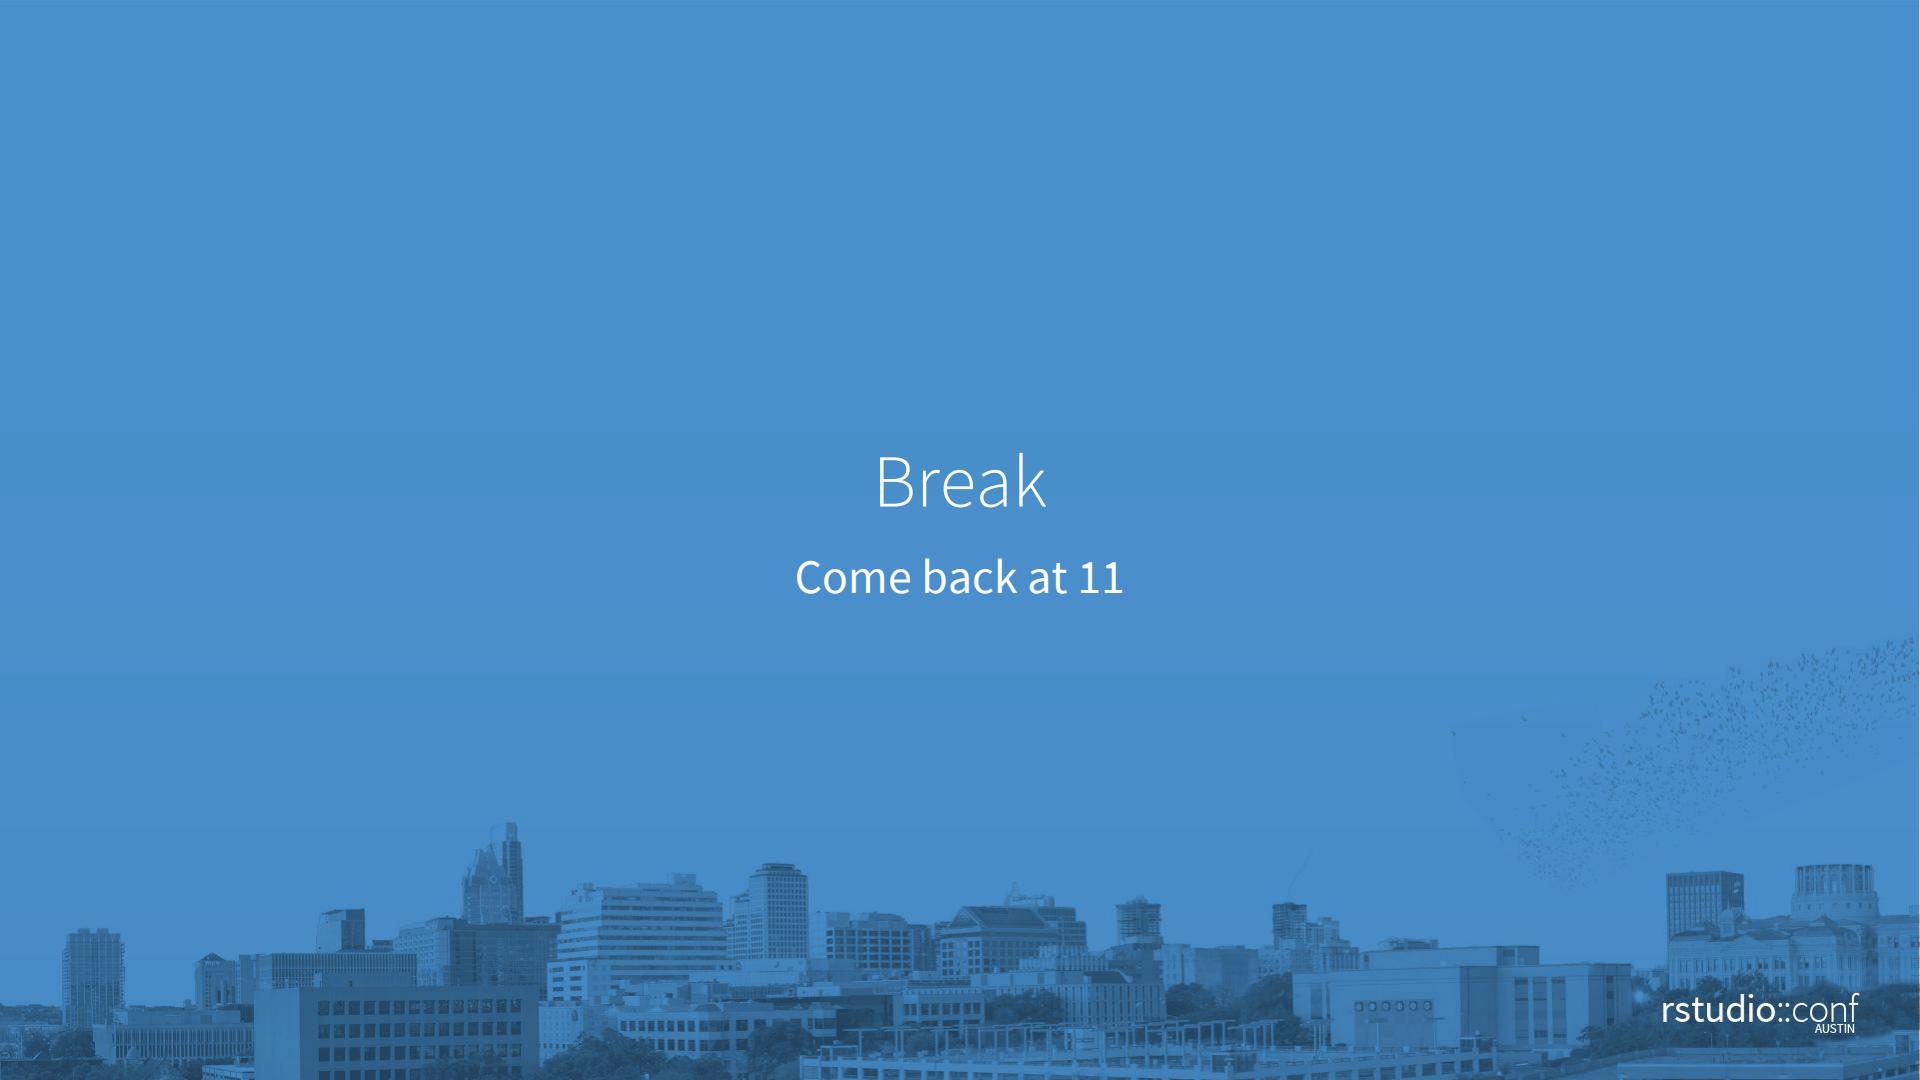
- First, what is up with the pre-run checklist? Any idea why these steps are (currently) necessary?
- The `runloadtest.R` file is going to start with a baseline test of 1 and then a test of 25. Why is the baseline important?

Deliverable:

Run the load test

Follow the first set of commands in `runloadtest.R`

How did the experience for 1 user compare to the experience for 25 users?

A dark blue-tinted photograph of the Austin, Texas skyline at dusk or night. The city lights are visible against a dark sky, with the Frost Bank Tower being a prominent feature on the left.

Break

Come back at 11

shinyloadtest: answers

What is up with our pre-run checklist?

```
# [] In use_models.R, un-comment the set.seed call on line 52
# [] Redeploy your app to RStudio Connect
# [] Make sure your app's access is set to "Anyone - no
login"
# [] Log out of RStudio Connect
```

} Recall: set.seed
} Bug being fixed

Compare the two runs. How did the experience for the 25 users compare to the 1 user?

```
loadtest <- load_runs(
  baseline = "answers/baseline",
  twentyfive = "answers/twentyfive"
)

shinyloadtest_report(loadtest)
```

Caching



renderCachedPlot - what?

Without:



With:



renderCachedPlot - how?

```
shinyOptions(cache = diskCache("./cache-dir"))
output$plot <- renderCachedPlot({
  plot(...)
}, cacheKeyExpr = list(input$student))
```

*specific to our example, normally use the object that represents a set of inputs that should generate a new output

renderCachedPlot 30 Minutes

First: Update your app code to use `renderCachedPlot`

Discussion:

caching + shinyloadtest

- How will introducing `renderCachedPlot` affect our load test experience?

Deliverable: *Re-run Load Test*

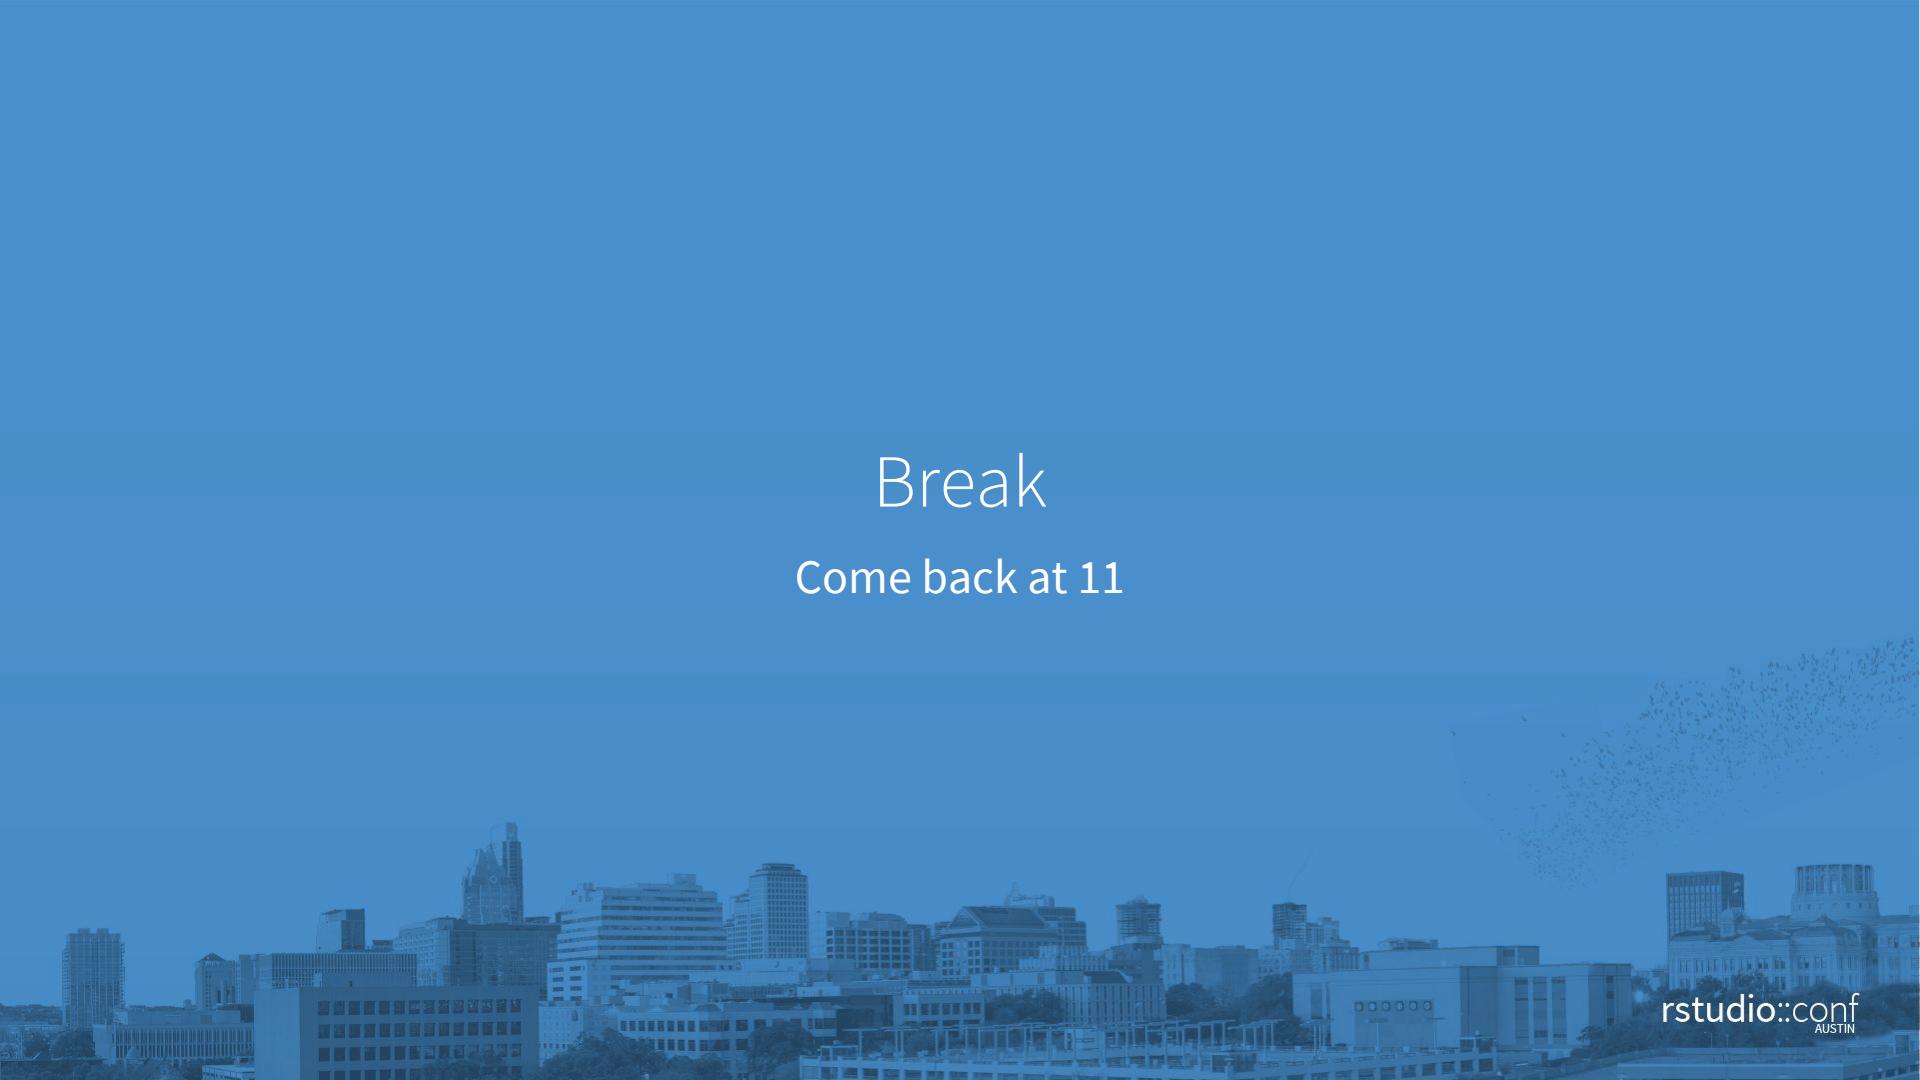
Redeploy the version of the app with `renderCachedPlot`

Re-run the load test and compare the outputs
(continue to follow along with `runloadtest.R`)

renderCachePlot: answers

answers/cached-app.R

```
loadtest <- load_runs(  
  baseline = "answers/baseline",  
  twentyfive = "answers/twentyfive",  
  cached = "answers/cached"  
)  
  
shinyloadtest_report(loadtest)d
```

A dark blue-tinted photograph of the Austin, Texas skyline at dusk or night. The city lights are visible against a dark sky, with the most prominent building being the Frost Bank Tower.

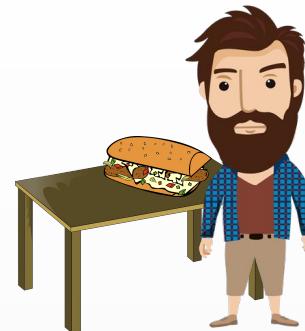
Break

Come back at 11

Scaling



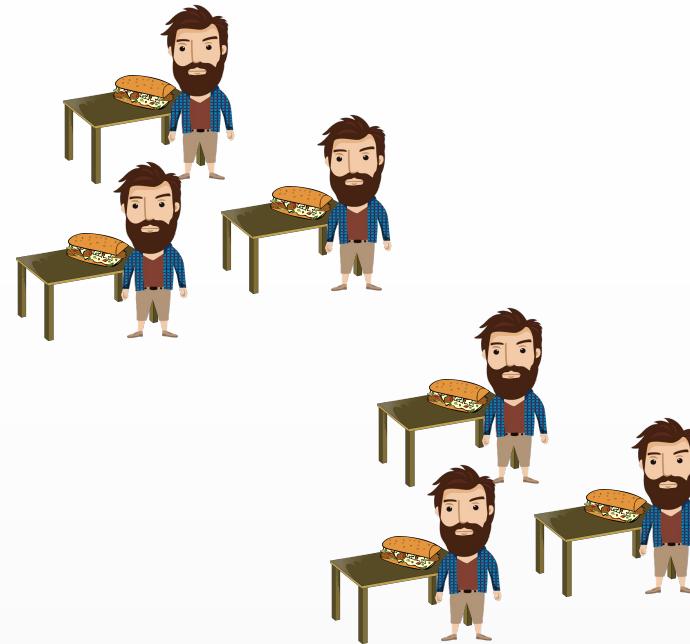
Scaling 101



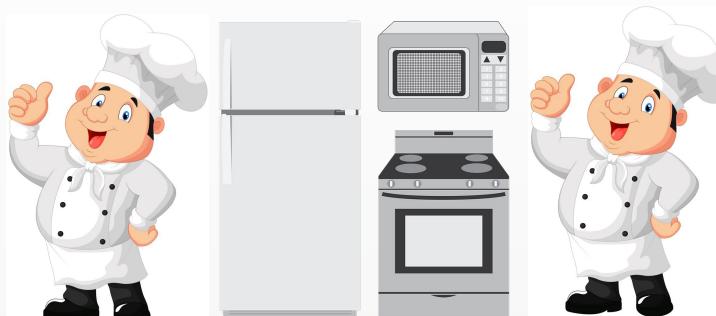
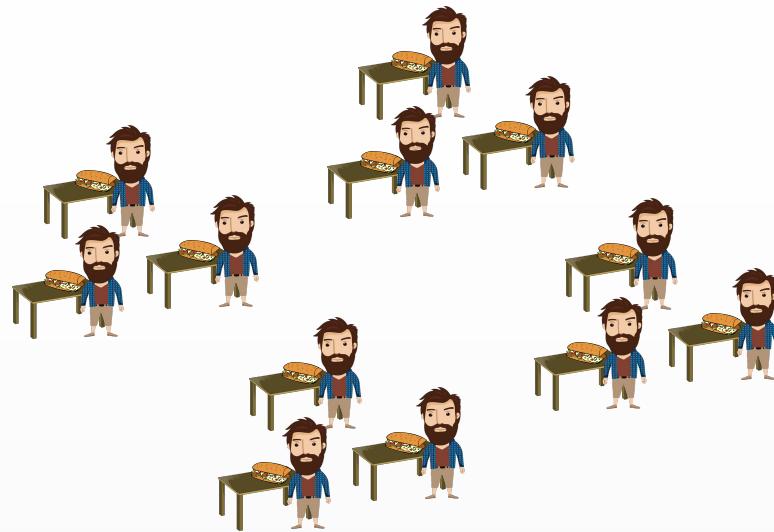
Scaling 101



Scaling 101



Scaling 101



Scaling 15 Minutes

First: Open the runtime settings for your app

Discussion:

Performance Tuning

- What do each of the settings mean?
- What changes to the settings might affect our app?

Deliverable: *Re-run Load Test*

Test your hypothesis by changing the runtime settings and then re-run the load test.

Compare the results to the prior test(continue to follow along with `runloadtest.R`)

Scaling: answers

```
loadtest <- load_runs(  
  baseline = "answers/baseline",  
  twentyfive = "answers/twentyfive",  
  cached = "answers/cached",  
  scaled = "answers/scaled"  
)  
  
shinyloadtest_report(loadtest)
```

Sharing is Caring

Metrics

Show history for months



CPU Using 0.03 of 4 cores



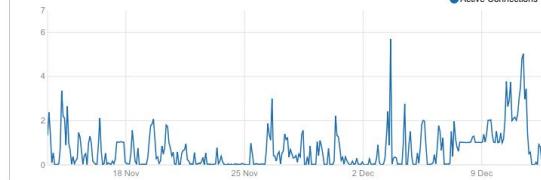
RAM Using 6.51GB of 14.7GB



Active Users 45 of 200 allowed



Shiny Connections 1 of 200 allowed



Processes

PID Content Name

5318 python_etl_twitter_app

5340 Access to Care Flex

Type

CPU

RAM

Shiny application

0.00

93.3 MB

Shiny application

0.00

58.4 MB

Sharing 10 Minutes

First: Open the admin dashboard in RStudio Connect

Discussion:

Multi-Tenant Environment

- What are the tradeoffs for modifying the scaling settings?
- What else might you need to consider when sharing infrastructure?
- How does these concerns change over time?

Deliverable: *Update checklist*

Create one pre-deployment check and one every-2-months check based on your discussion.

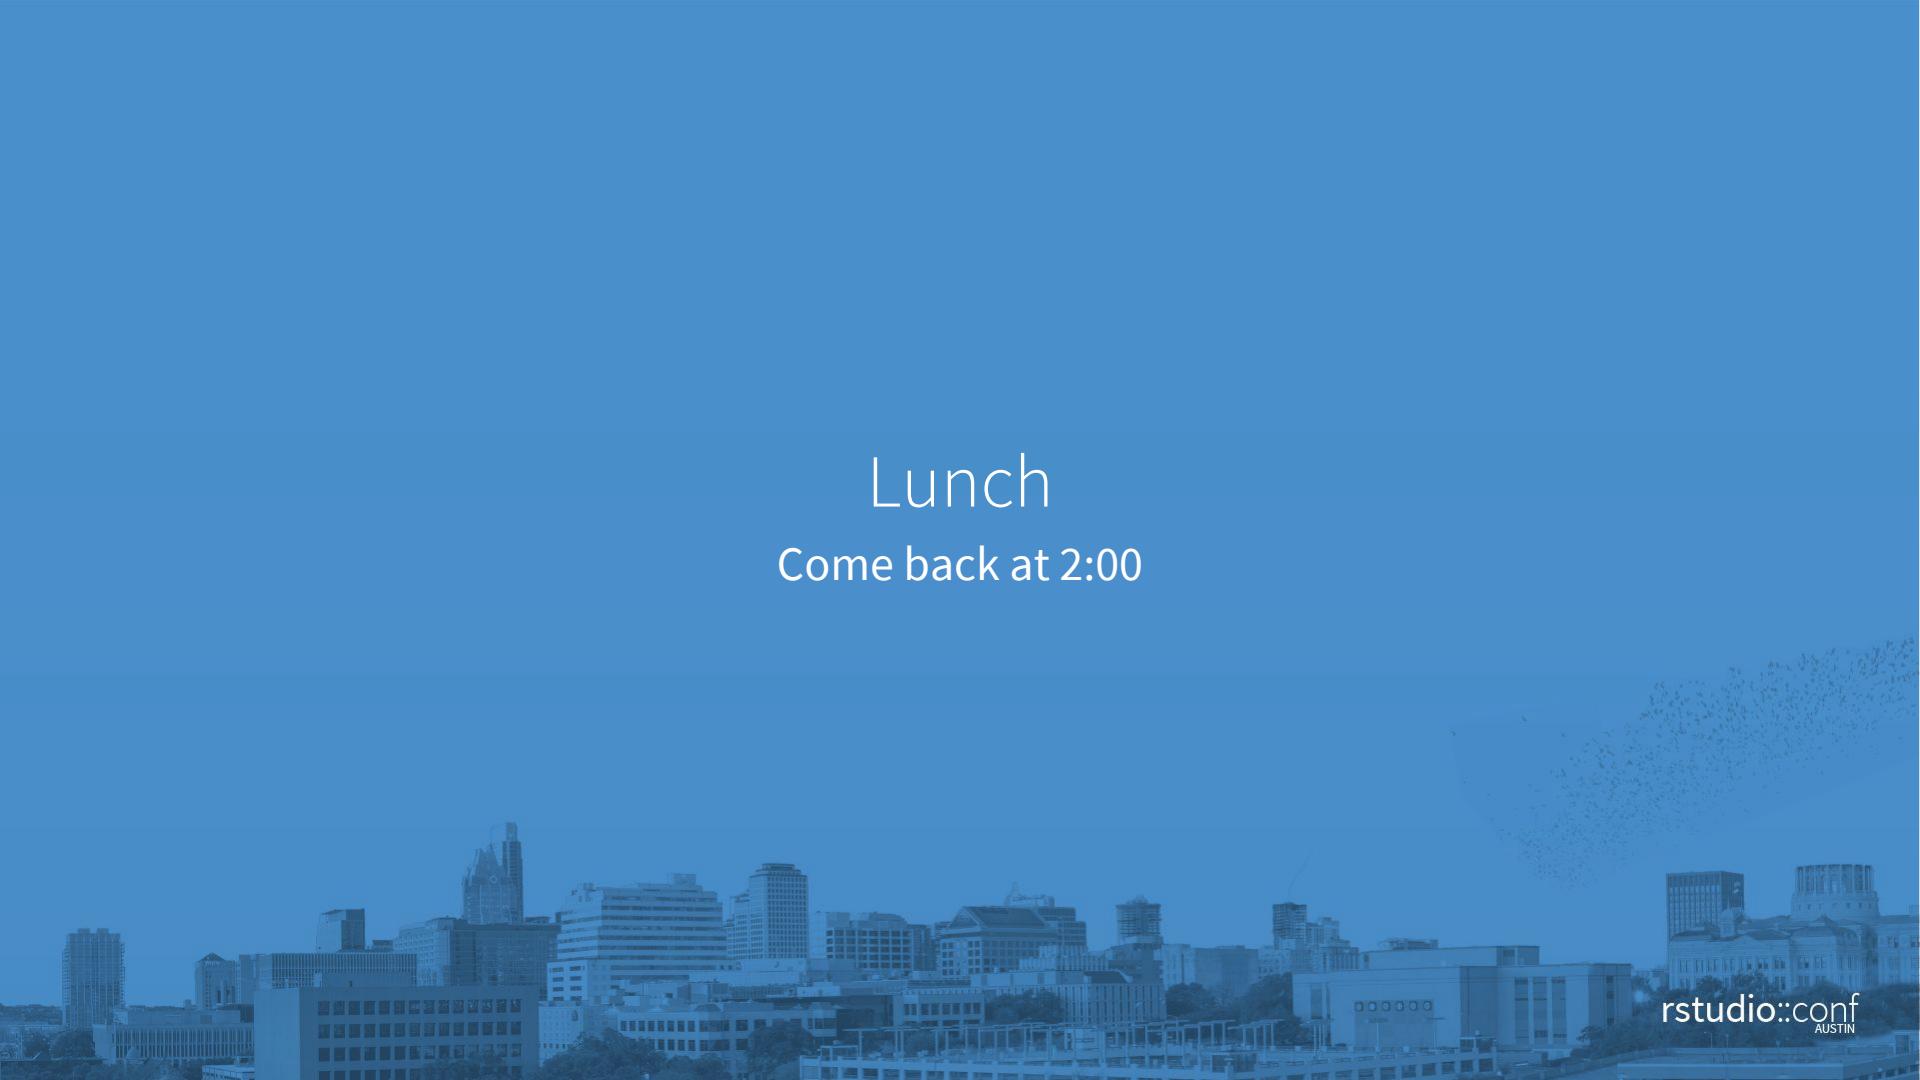
Sharing is Caring: answers

Pre-Deployment:

[] What other content is running on the server? Do I have enough resources for my app?

Every 2 Months:

[] Check logs to ensure server resource consumption is reasonable. Can I automate this?

A dark blue-tinted photograph of the Austin, Texas skyline at dusk or night. The city lights are visible against a dark sky, with the tallest building, the Frost Bank Tower, clearly visible on the left.

Lunch

Come back at 2:00

Alternatives to Shiny

Review survey results for plumber and R Markdown familiarity

R Markdown

Scheduled data updates, email distribution, and client side interactivity.



Browser and R are connected!



HTML gets rendered. Time passes.
Then the server gives the HTML to
the browser.

R Markdown 15 Minutes

First: Create your own R markdown document (New File)

Discussion:

Consider R Markdown

- How does an R Markdown document scale?
- What types of interactivity can be added to R Markdown?

Deliverable: *Create a decision matrix comparing Shiny to R Markdown*

<i>Alternatives Criteria</i>	Shiny	R Markdown
Criteria A		
Criteria B		
Criteria C		
Criteria D		

R Markdown: answers

<i>Criteria</i>	<i>Alternatives</i>	R Markdown
# of Inputs	Users need multiple levels of dependent inputs to get an answer	All inputs can be specified up front
Update Frequency	Updates need to happen in real time	Updates occur on a regular basis
Output	Users come to the system to do exploratory Q&A, the outputs are the results	Users want a copy of the results or want results distributed to their email
Interactivity	Interactivity depends on lots of data	Interactivity can be based on data sent to the browser

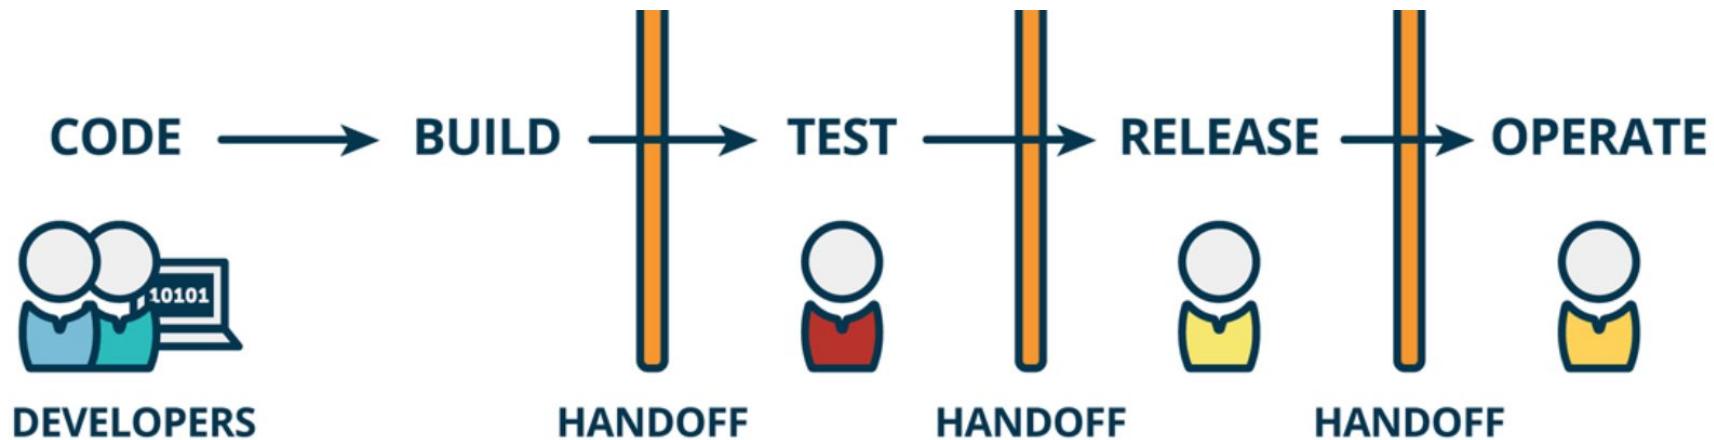
Adventures in Mixing Shiny Apps with IT Groups

DevOps 101



Software Developers Work with IT and so can YOU

Managing Code Deployment Handoffs



mindtheproduct.com

It's not “your job” to understand operations and systems administration.

It's not “their job” to understand R programming. SO HOW DO WE MAKE THIS LESS PAINFUL??

Steal Existing Strategies for Managing Code Handoffs & Define Shared Goals

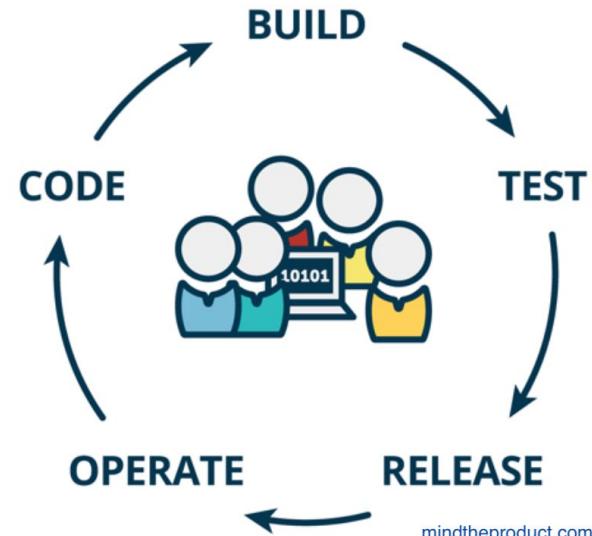
DevOps: Self-help philosophy for IT and software development teams

Solve a core problem:

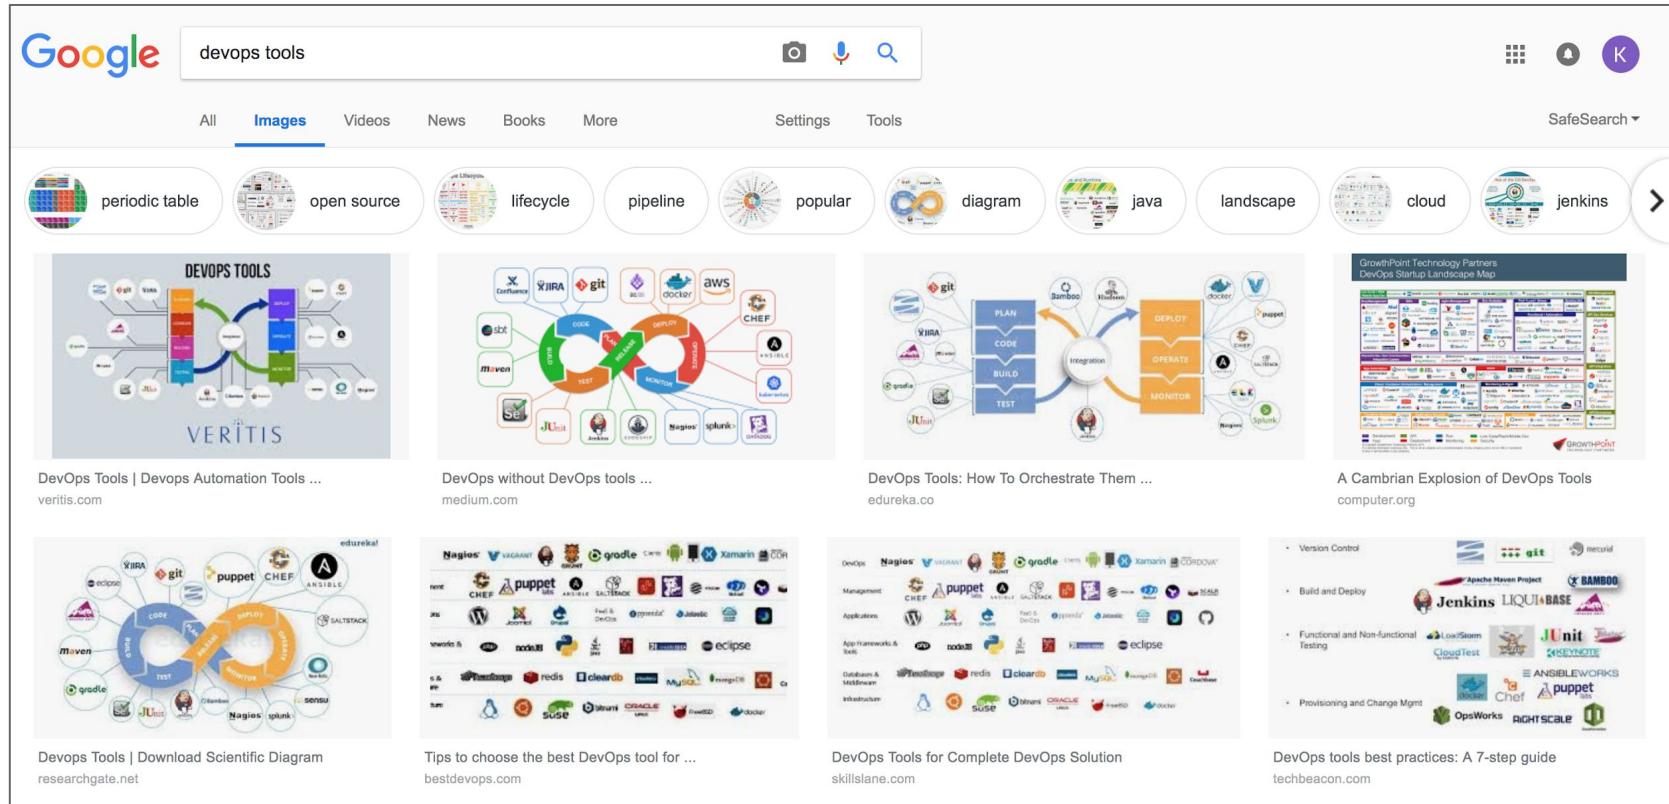
Everyone fears taking code into production

Code deployments are HIGH RISK EVENTS (performed rarely)

And they need to be TRIVIAL NON-EVENTS (daily work)



DevOps - Land of 1000 Intimidating New Tools

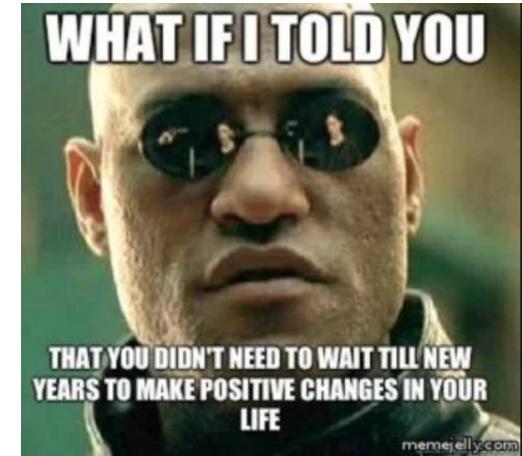


1. DevOps is a philosophy or set of best practices
2. It's a new approach to software development and delivery
3. **There's nothing new in DevOps**

A framework for making sense out of common sense

Are DevOps Principles just nerdy New Year's Resolutions?

- **Flow:** Accelerate the delivery of work from development to operations to customers
- **Feedback:** Create ever safer systems of work
- **Continual Learning and Experimentation:** A high-trust culture and scientific approach to organizational improvement risk-taking as part of daily work



-- Support --

Technology Value Stream: The process required to convert a business hypothesis into a technology-enabled service that delivers value to the customer

The DevOps Handbook

Three principles form the underpinnings of DevOps:

1. Accelerate Flow

- Make work visible
- Limit Work in Progress (WIP)
- Reduce Batch Sizes
- Reduce the number of handoffs
- Continually identify and elevate constraints
- Eliminate hardships and waste

2. Utilize Feedback

- See problems as they occur
- Swarm to solve problems and build new knowledge
- Keep pushing quality closer to the source
- Enable optimizing for downstream work centers

3. Learn and Experiment

- Enable organizational learning and a safety culture
- Institutionalize the improvement of daily work
- Transform local discoveries into global improvements
- Inject resilience patterns into daily work

Principles are business-y and it takes some digging to get to the actionable self-help advice.



The DevOps Handbook:
How to Create World-Class
Agility, Reliability, and...
Gene Kim

Everyone's Production is
Different

And Evolves at a Different Rate

Production is...

DOCUMENTATION - TESTING & MONITORING

- Creating apps that can reach a wider audience and are deployed/tested in a consistent manner
- Running in a way that is stable to use, documented and monitored

AT SCALE

- Scaled to a larger audience
- Bulletproof, scalable, fails predictably
- Live to 1000s of users with production vehicle data

Credibility

CUSTOMER/USER FACING

- Ready to use
- Software that end users are using
- An app that is live and available to the end user
- Apps on our production server are available to our clients
- Client facing

ENVIRONMENTAL REQUIREMENTS

- An area where validated applications are deployed in a locked down environment
- The main part of a company that handles all process
- Application or system operates effectively without much maintaining effects
- A server or environment that runs the “final” applications that your ultimate end-users (often external customers) use to get stuff done

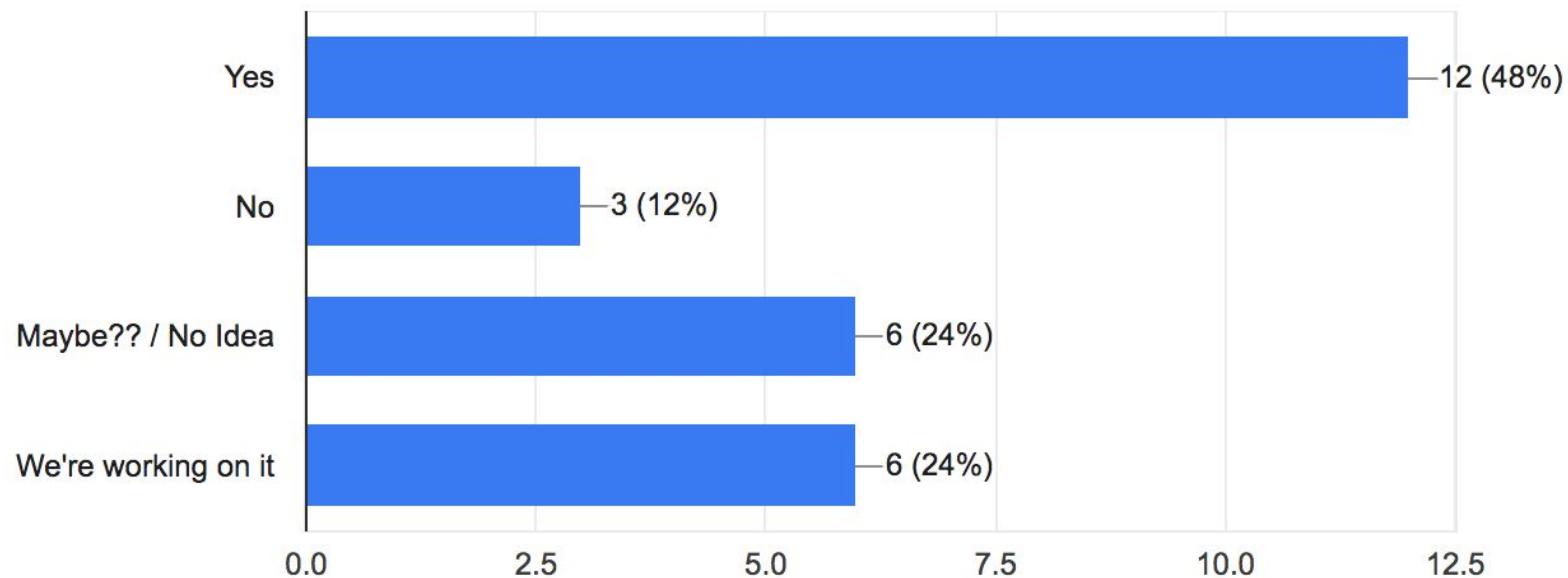
SERVICE LEVEL AGREEMENTS

- Required for mission-critical operations; downtime affects the ability to serve customers
- Deployed for end users to have continual access without performance issues

My organization has a team or process in place for handling code deployment



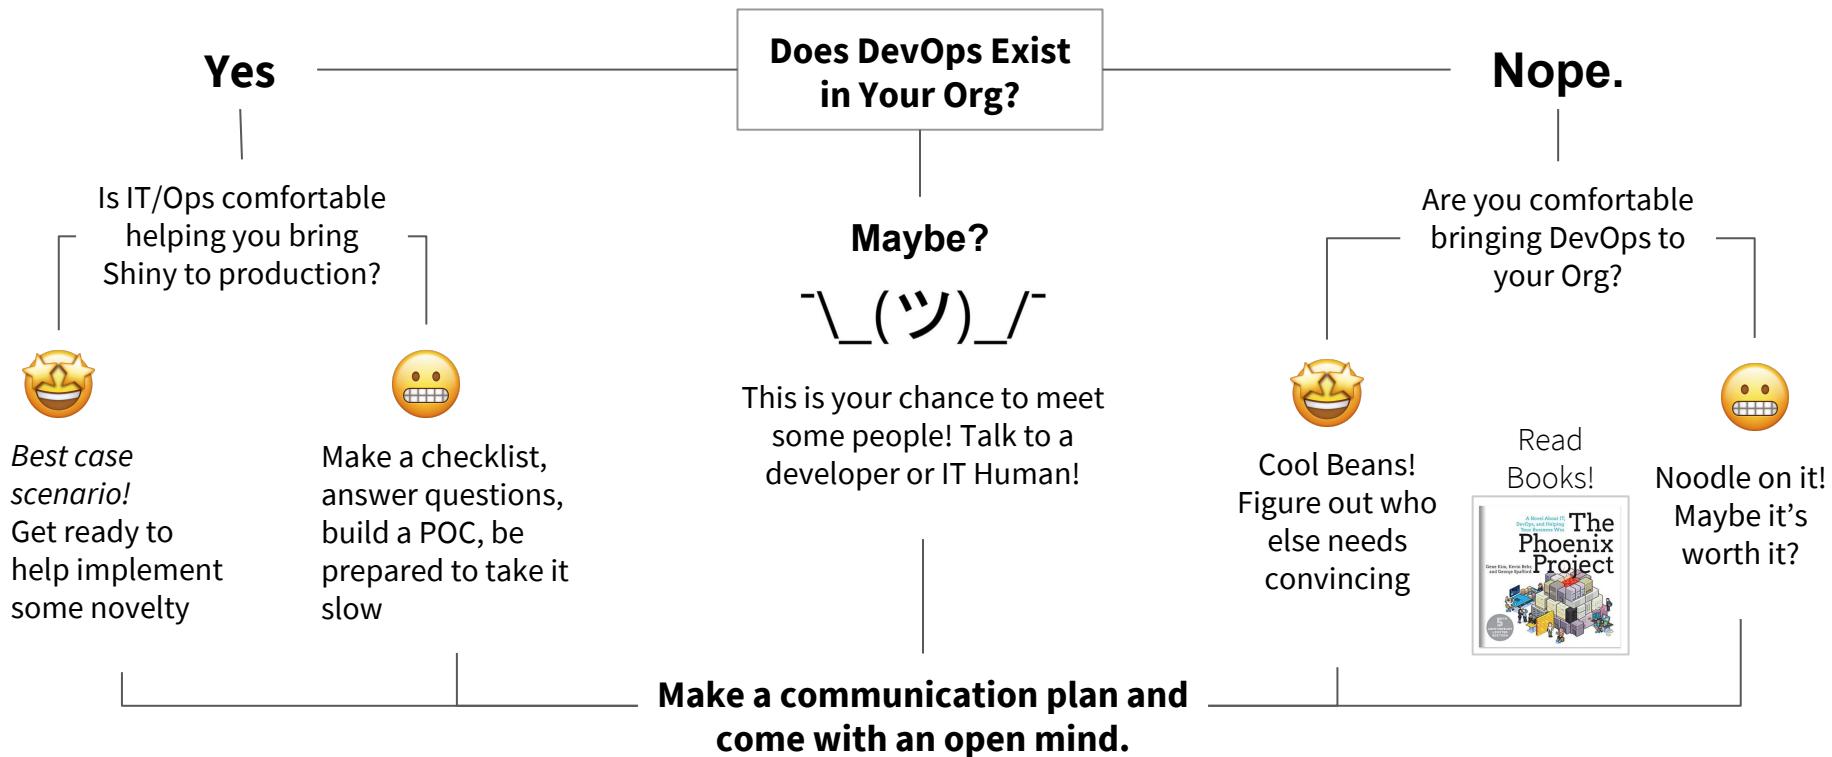
25 responses



Put a name on what you're working with

Pathological (power-oriented)	Bureaucratic (rule-oriented)	Generative (performance-oriented)
Low cooperation	Modest cooperation	High cooperation
Messengers shot	Messengers neglected	Messengers trained
Responsibilities shirked	Narrow responsibilities	Risks are shared
Bridging discouraged	Bridging tolerated	Bridging encouraged
Failure leads to scapegoating	Failure leads to justice	Failure leads to enquiry
Novelty crushed	Novelty leads to problems	Novelty implemented

How to wade in ... with Empathy and Strategery!



Empathetic Communication is Challenging

Start by answering some questions...

- What is a Shiny Application?
- Who is the audience?
- What is your service level agreement definition? (SLA)
- What does your analytic architecture look like today?
 - What are your goals for evolving this architecture?
- How will monitoring be handled?
- Who is responsible for maintenance?

Make work visible - Build a comprehensive checklist - Optimize

Don't get caught up in the hype!



There are no easy answers

The search results page displays several images related to DevOps tools:

- A diagram titled "DEVOPS TOOLS" by VERITIS, showing a complex network of interconnected tools.
- A diagram titled "DevOps without DevOps tools ..." by medium.com, showing a circular flow of processes.
- A diagram titled "DevOps Tools: How To Orchestrate Them ..." by edureka.co, showing a flowchart of tool integration.
- A diagram titled "Devops Tools | Download Scientific Diagram" by researchgate.net, showing a circular diagram of various DevOps tools.
- A diagram titled "Tips to choose the best DevOps tool for ..." by bestdevops.com, showing a grid of tool logos.
- A diagram titled "DevOps Tools for Complete DevOps Solution" by skillslane.com, showing a grid of tool logos.
- A diagram titled "A Cambrian Explosion of DevOps Tools" by computer.org, showing a dense grid of tool logos.

There is no “typical workflow diagram”

The *standard architecture* might not make sense for you

There is no perfect deployment pipeline - everything is *and* should be evolving

*Even more important than daily work is
the improvement of daily work*

-- Mike Orzen, Lean IT





Allison Moberger 🍔

@allimoberger



"You shouldn't feel ashamed about your code - if it solves the problem, it's perfect just the way it is. But also, it could always be better."

-@hadleywickham #rstudioconf

2:12 PM · 15 Jan 19 · Twitter Web Client

69 Retweets 220 Likes



Allison Moberger 
@allimoberger

EDIT:

"You shouldn't feel ashamed about
your analytic infrastructure and code - if it solves the problem,
deployment patterns
it's perfect just the way it is. But
also, it could always be better."
-[@hadleywickham](#) #rstudioconf

2:12 PM · 15 Jan 19 · [Twitter Web Client](#)

69 Retweets 220 Likes

DevOps Foundations: Start Setting Goals

Supplement Book: bit.ly/shiny-prod

Goal: Develop in Production-like Environments

What do your dev, test and production environments look like?

- Do developers have access to production-like environments on their own workstations?

- Can these environments be assembled on-demand?

- Version Control



- Infrastructure as Code



Exercise: Relate it Back to R and Shiny

1. In what ways can analytic infrastructure for R change between development and production environments?
2. What processes or tools could help solve these issues?
3. *Should development inform how a production environment is built, or should production inform how an application is developed?*

Goal: Make Code Deployment a Low-Risk Process

Why should data scientists design deployment processes?

You bring a uniquely informed perspective.

Having developers focus on automating and optimizing the deployment process can lead to significant improvements in deployment flow.

- Smoke testing
- Build validation tools
- Environment consistency



Exercise: Relate your Shiny deployment checklist to the pipeline an IT group might consider. Where is there risk? How can it be reduced?

- Packaging code in ways suitable for deployment
- Creating pre-configured virtual machine images or containers
- Automating the deployment and configuration of middleware
- Copying packages or files onto production servers
- Restarting server, applications or services
- Generating configuration files from templates
- Running automated smoke tests to make sure the system is working and correctly configured
- Running testing procedures
- Scripting and automating database migrations

Goal: Decouple Deployments from Releases

Deployment is any push of code to an environment (test, prod)

Release is when that code (feature) is made available to users or customers

Deployment on demand and thoughtful release strategies allow more control (and more success) over the delivery of features to end users.

Environment-based Release Patterns

- Blue-Green Deployment Pattern
- Canary and Cluster Immune System

Application-based Release Patterns

- Feature Toggles
- Dark Launches

Exercise: Which of these release patterns is most well-suited for Shiny Application development and deployment with the RStudio Connect publishing platform?

Case Studies 30 Minutes

- A) Before publishing an update to our app, we want to run user acceptance tests.
- B) All of our production code must be deployed from Git.
- C) We scale applications through Docker - how does Shiny fit in?

First: Pick one case study

Discuss:

What tools does your admin team use?

How would you explain the goals of a shiny app to DevOps?

Deliverable

Whiteboard (in a medium of your choice) what an architecture might look like for your case study.

Delineate who (and what tool) is responsible for code, R, R packages.

Break

Come back at 4:00

Case Studies 10 Minutes

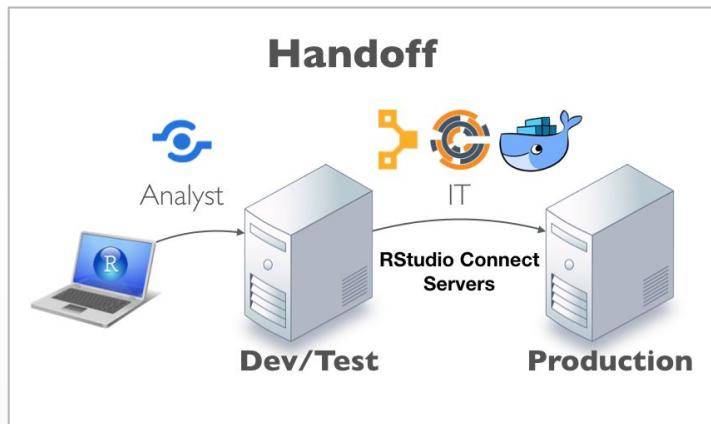
Re-acquaint yourself with the case study you picked.

Case Study A

Before publishing an update to our app, we want to run user acceptance tests.

Case Study A answers

Before publishing an update to our app, we want to run user acceptance tests.



Code: Committed to Git by users after being deployed to Dev.

R: Multiple versions installed on server by admin's Ansible scripts.

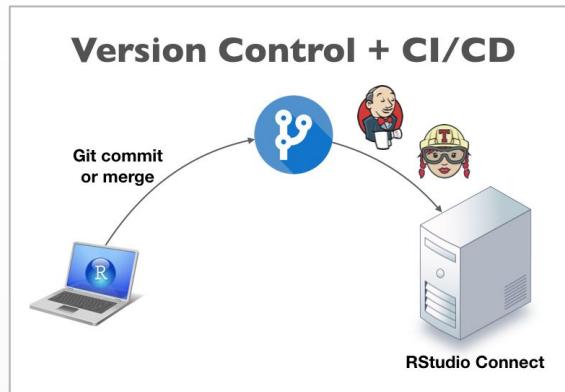
R packages: Restored by RStudio Connect

Case Study B

All of our production code must be deployed from Git.

Case Study B answers

All of our production code must be deployed from Git.



Code: Committed to Git by users.
Deployed by Jenkins.

R: Multiple versions installed on server by admin's Ansible scripts.

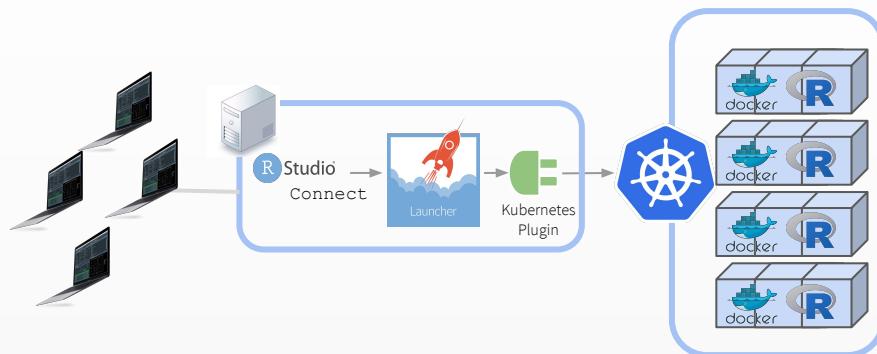
R packages: Restored from a manifest in Git during Jenkins runtime.

Case Study C

We scale applications through Docker - how does Shiny fit in?

Case Study C answers

We scale applications through Docker - how does Shiny fit in?

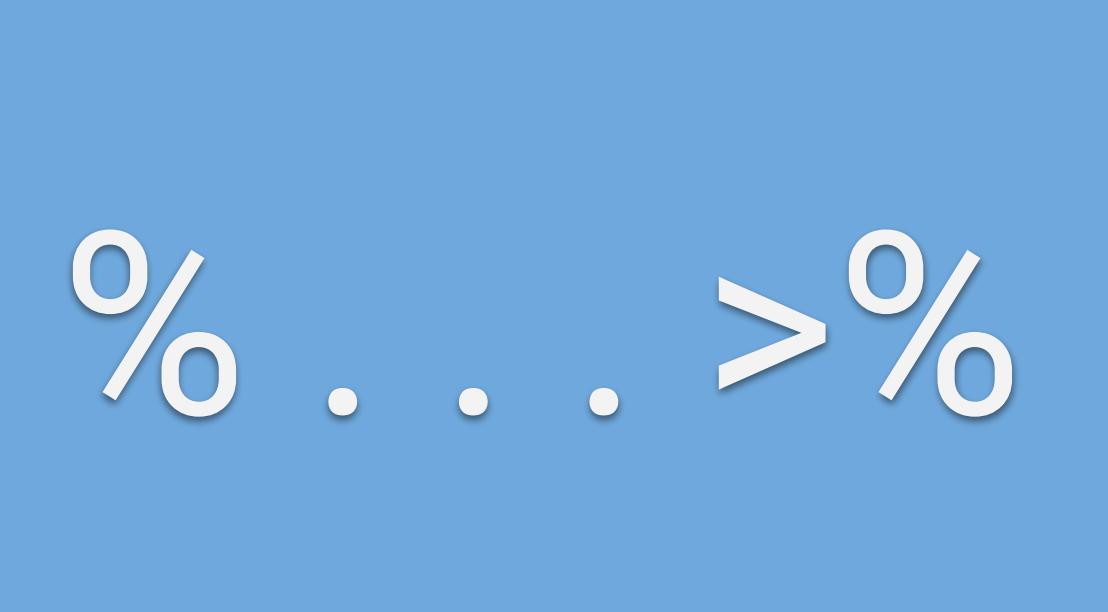


Code: Cloned into GitHub as a step in the Dockerfile.

R: Inherited from a base image.

R packages: Installed as a step in the Dockerfile.

Async



% . . . >%

[Checkout the Webinar](#)

Wrap Up

What does our production-ready checklist look like?

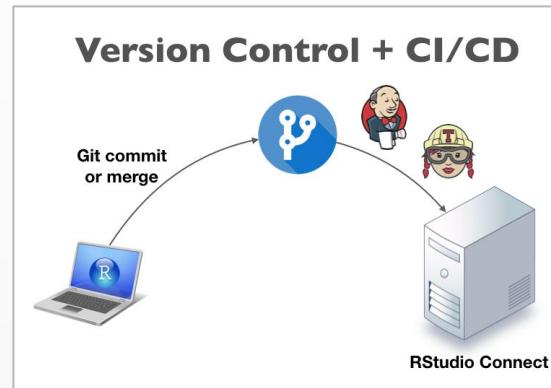
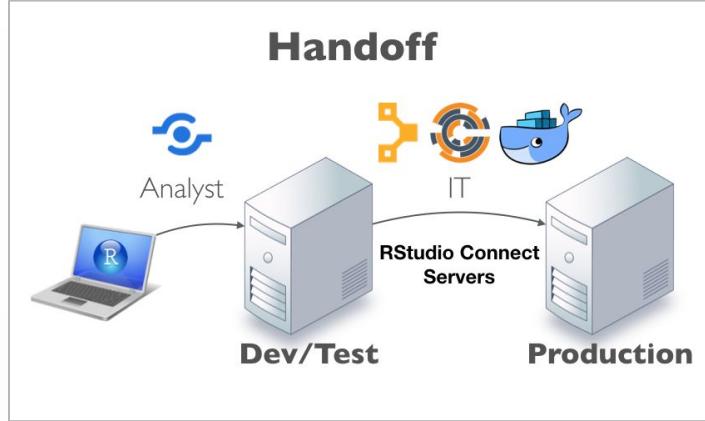
GitHub Material >> bit.ly/shiny-prod-class

Supplement Book >> bit.ly/shiny-prod-book

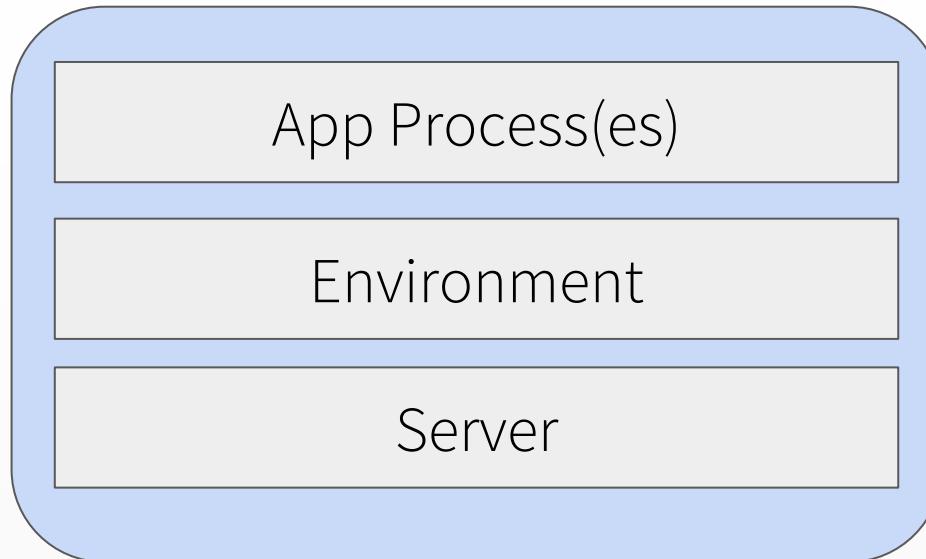
There's a survey! >> <https://bit.ly/2VFQRPo>

Appendix (aka Garbage Slides)

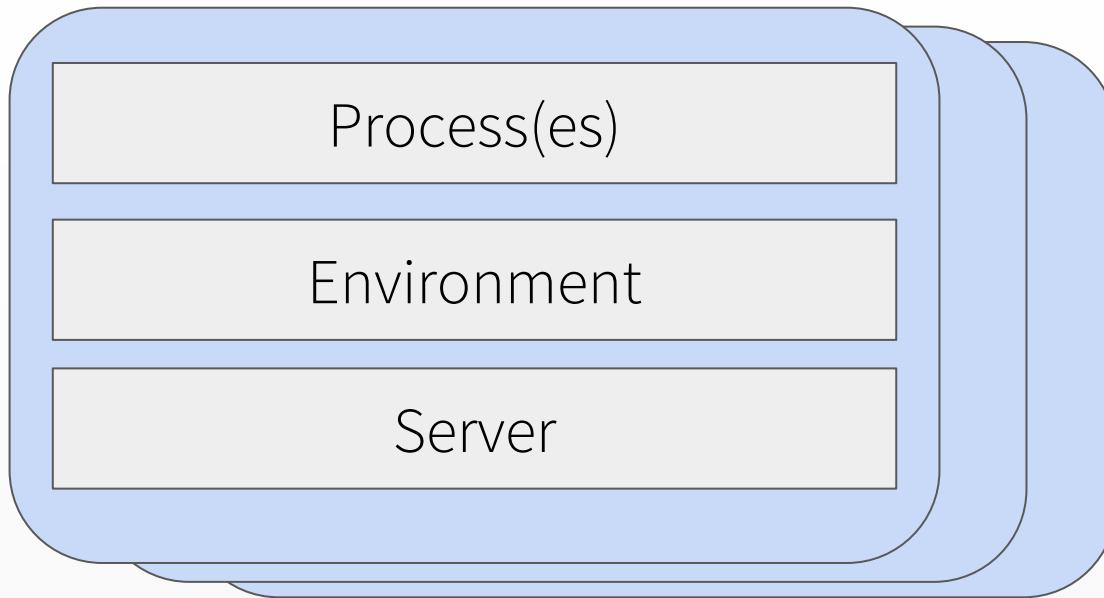
DevOps



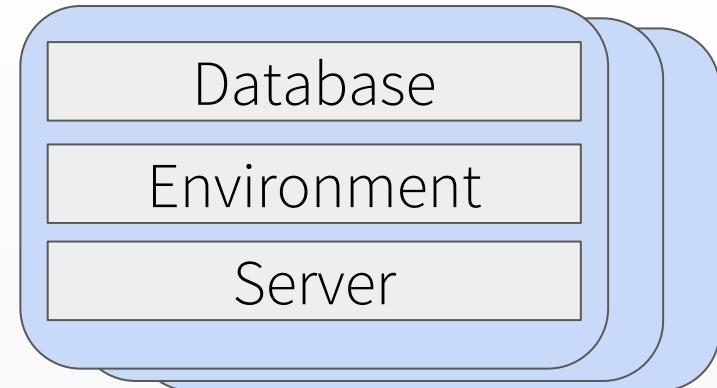
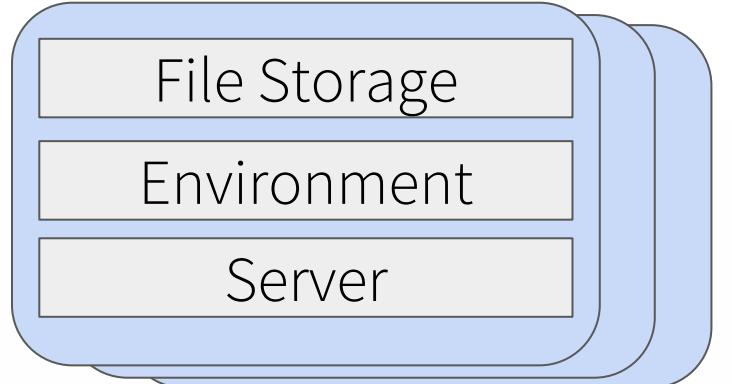
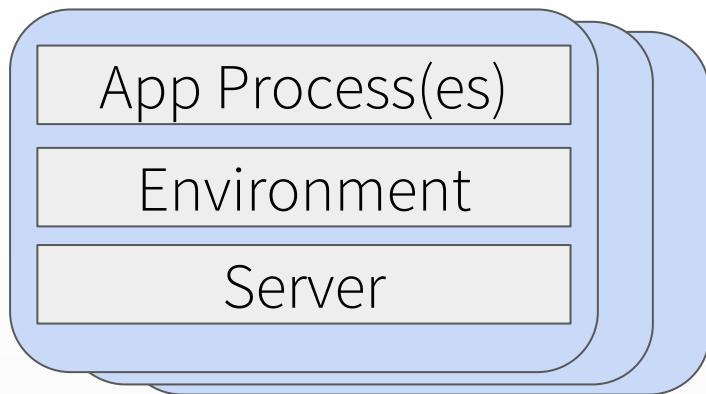
DevOps Vocab 101



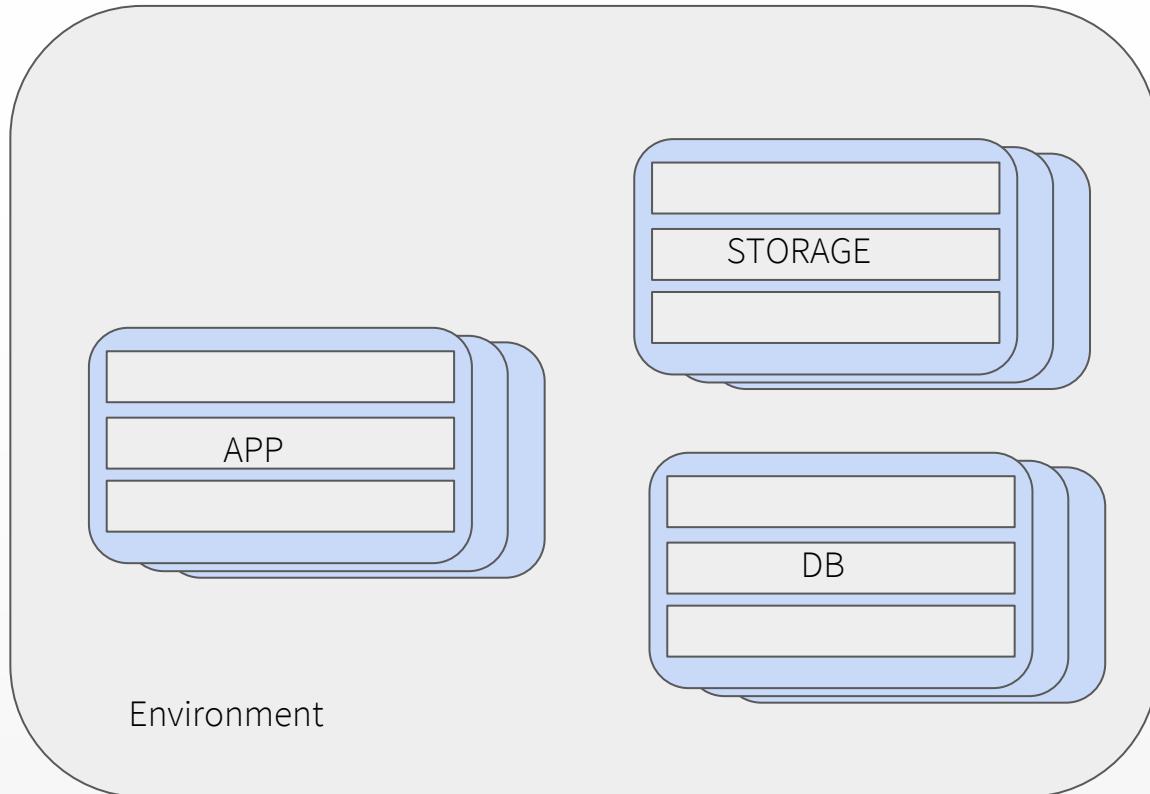
DevOps Vocab 101



DevOps Vocab 101



DevOps Vocab 101



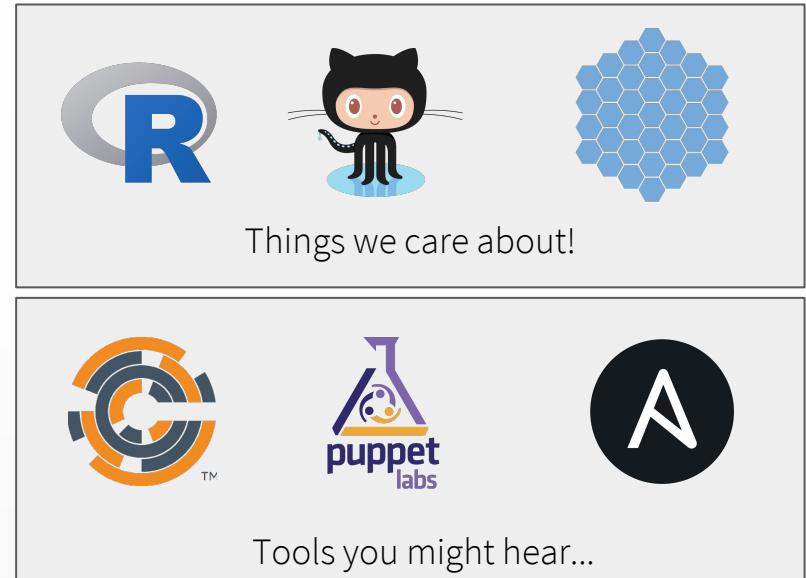
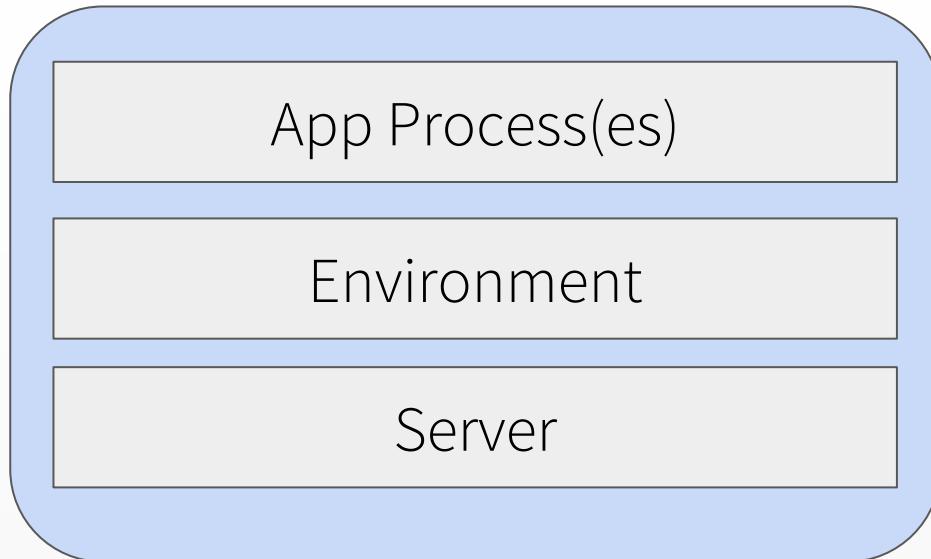
DevOps Vocab 101

DEV

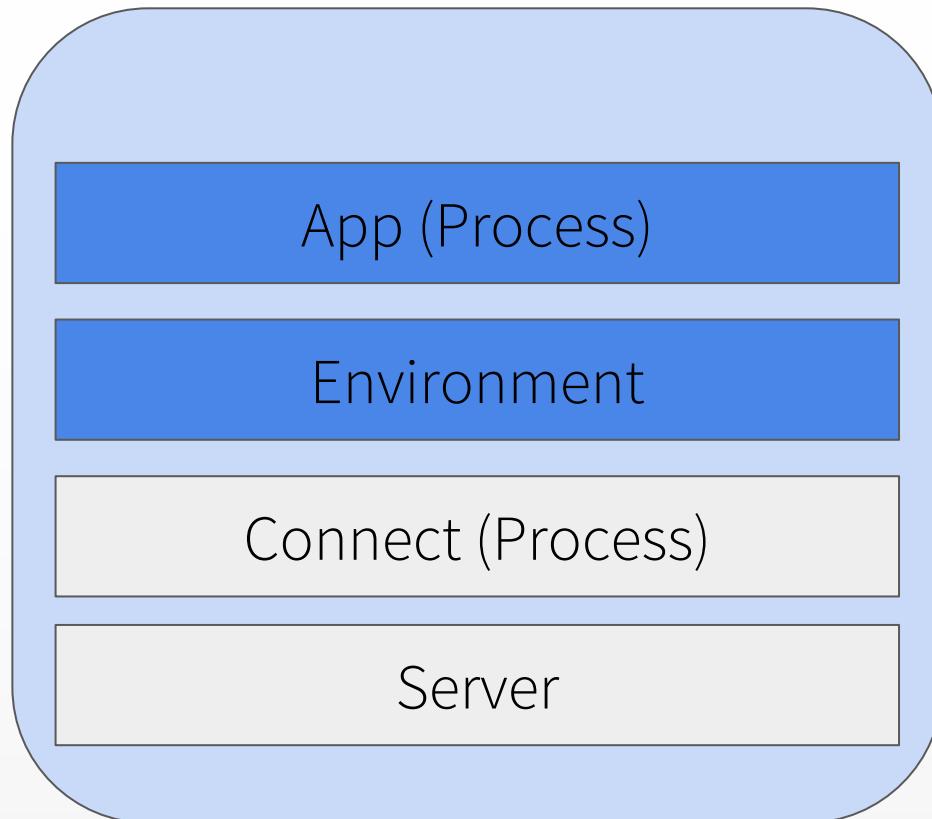
QA

PROD

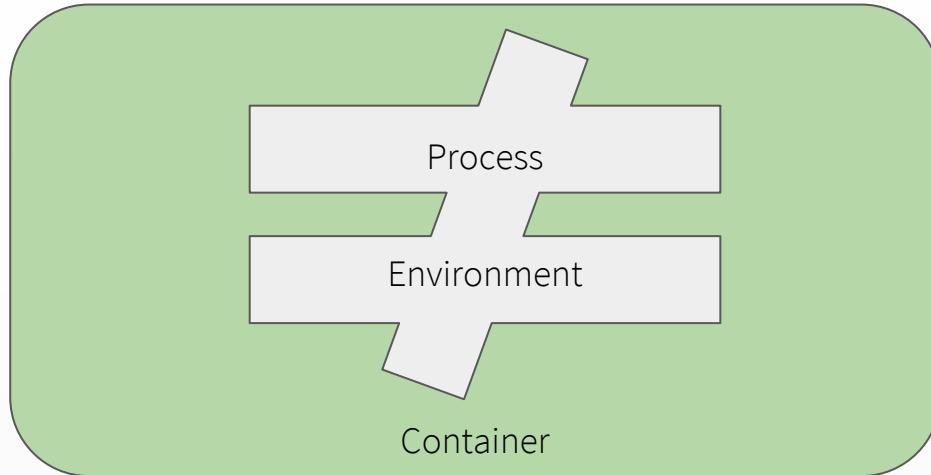
DevOps Vocab 101



DevOps Vocab 101 - The RStudio Connect Version

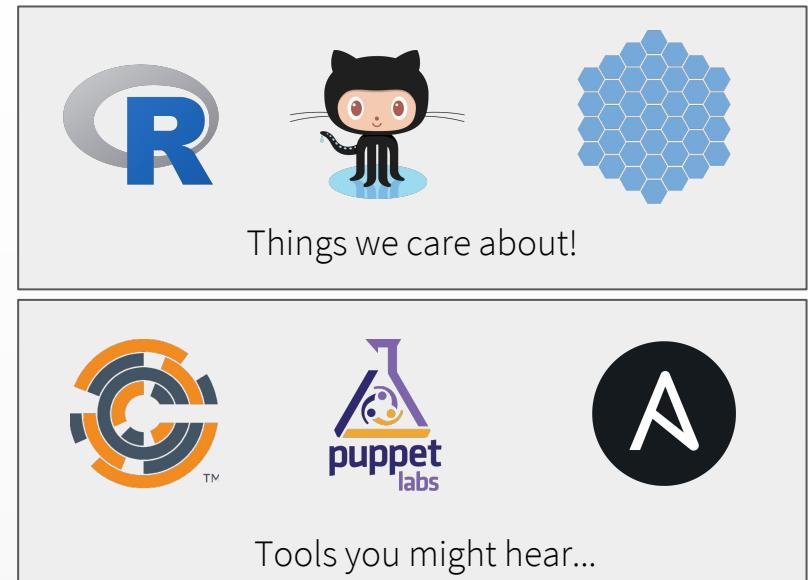
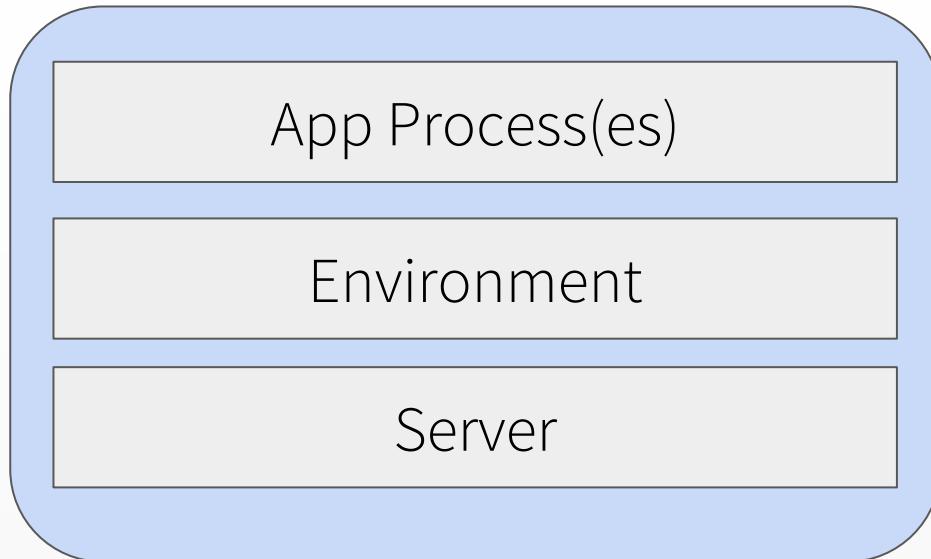


DevOps Vocab 101 - The Docker Version

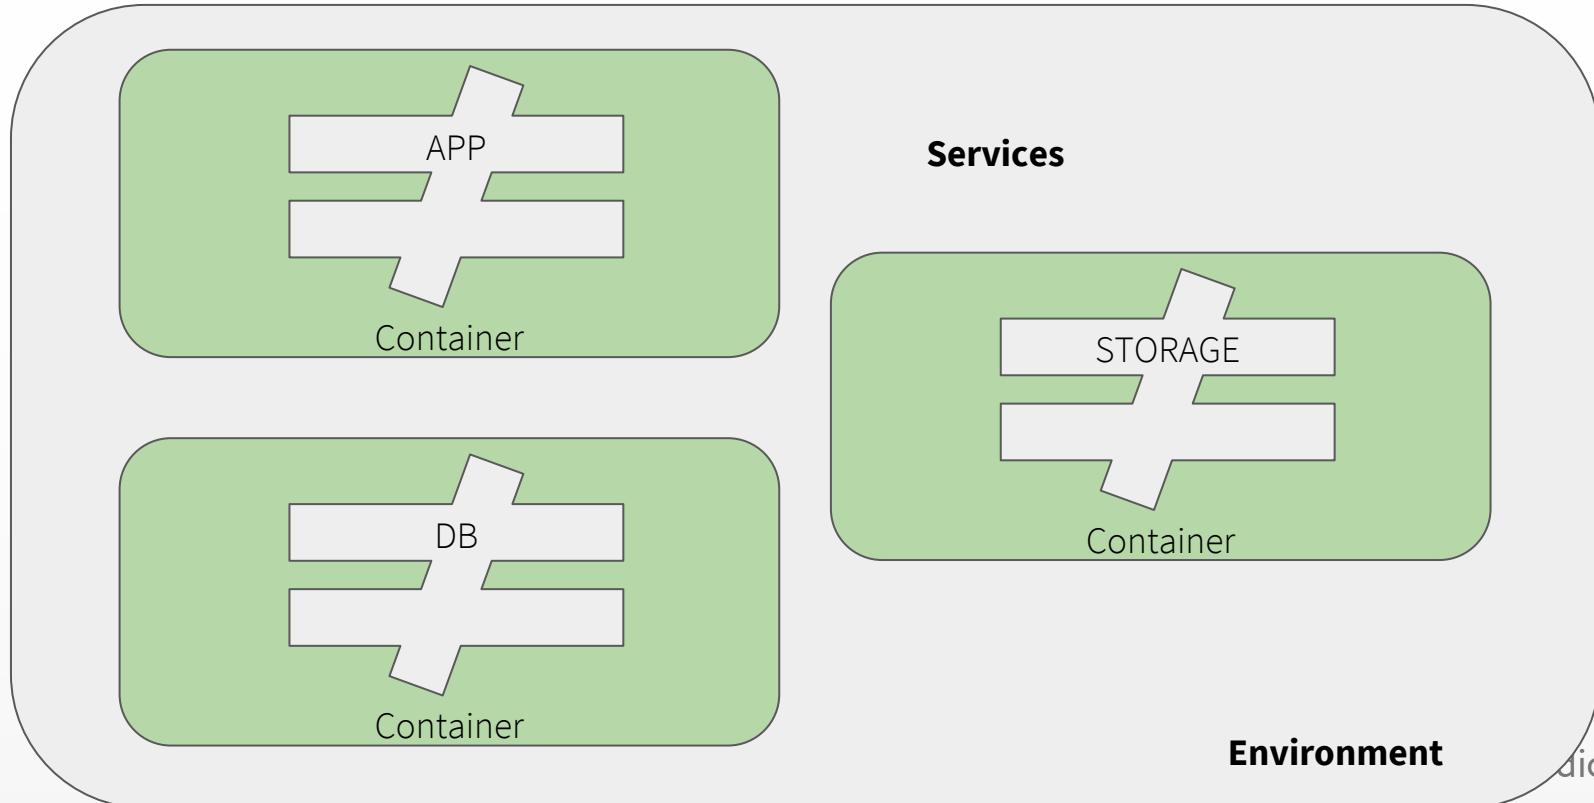


“ Float like an environment, sting like a process.

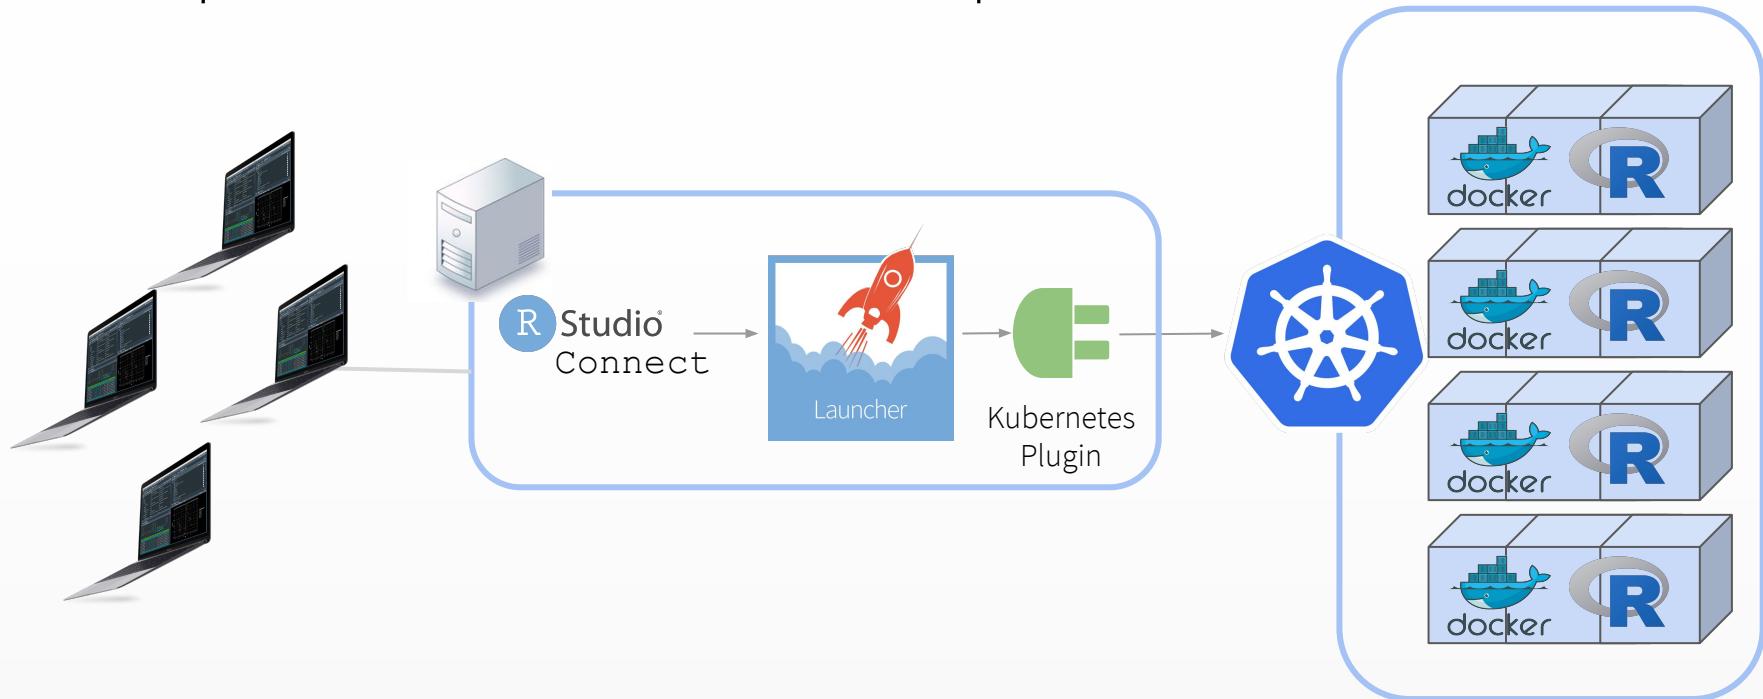
TBD



DevOps Vocab 101 - Kubernetes?



DevOps Vocab 101 - The Mashup



DevOps Vocab 102 - Other Hype Concepts



Serverless
architecture



High-availability

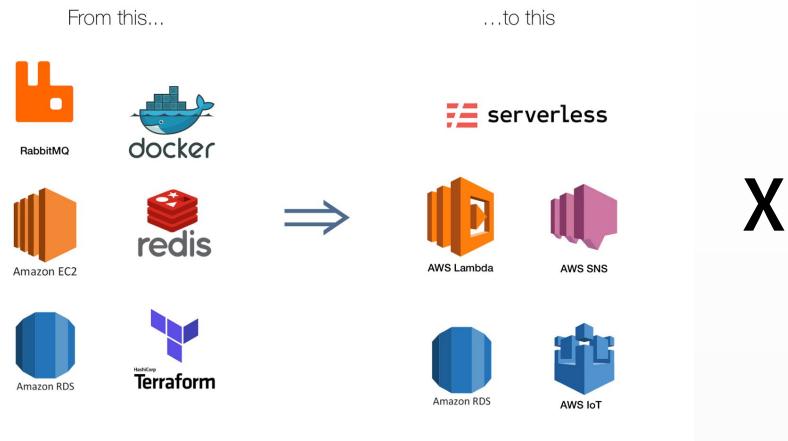


Event-driven



Zero administration

DevOps Vocab 102 - Names for Things



X

